



deeplearning.ai

Generative Models


Outline

- What are generative models
- Types of generative models



Generative Models vs. Discriminative Models

Discriminative models



Features Class

$$X \rightarrow Y$$

$$P(Y|X)$$

Generative models



Noise Class Features

$$\xi, Y \rightarrow X$$

$$P(X|Y)$$

Generative Models vs. Discriminative Models



Generative models



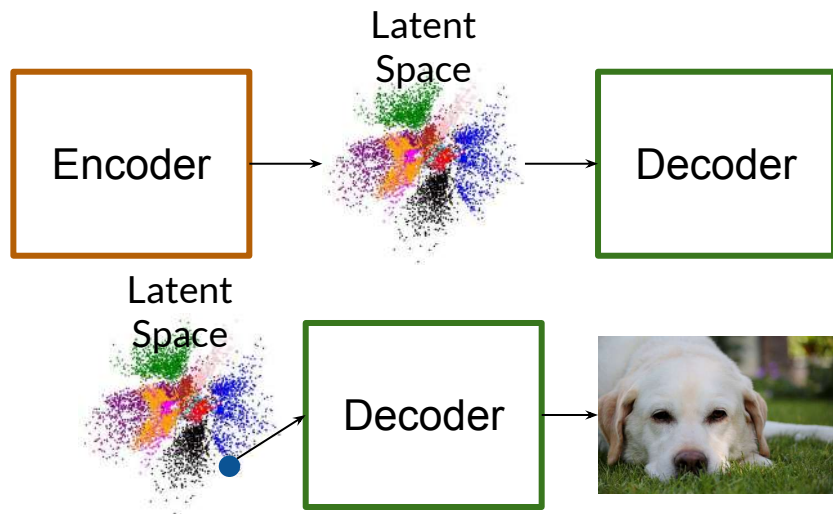
Noise Class Features

$$\xi, Y \rightarrow X$$

$$P(X|Y)$$

Generative Models

Variational Autoencoders




Generative Adversarial Networks


Available from: <https://arxiv.org/abs/1804.00891>

Generative Models

Variational Autoencoders




Generative Adversarial Networks



Available from: <https://arxiv.org/abs/1804.00891>

Summary

- Generative models learn to produce examples
- Discriminative models distinguish between classes
- Up next, GANs!






deeplearning.ai

Real Life GANs

Outline

- Cool applications of GANs
- Major companies using them



GANs Over Time



Ian Goodfellow
@goodfellow_ian




4.5 years of GAN progress on face generation.

arxiv.org/abs/1406.2661 arxiv.org/abs/1511.06434

arxiv.org/abs/1606.07536 arxiv.org/abs/1710.10196

arxiv.org/abs/1812.04948



GANs Over Time



Face Generation
StyleGAN2

These people do
not exist!

Karras, Tero, et al. "Analyzing and improving the image quality of stylegan." *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2020.

GANs Over Time



StyleGAN2



Mimics the
distribution of the
training data

Karras, Tero, et al. "Analyzing and improving the image quality of stylegan." *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2020.

<https://9gag.com/gag/aWYZKWx>

GANs for Image Translation


From one domain to another

CycleGAN



Park, Taesung, et al. "Semantic image synthesis with spatially-adaptive normalization." *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*. 2019.

GANs for Image Translation



GauGAN

Doodles



Pictures

Park, Taesung, et al. "Semantic image synthesis with spatially-adaptive normalization." *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*. 2019.

GANs are Magic!



Zakharov, Egor, et al. "Few-shot adversarial learning of realistic neural talking head models." *Proceedings of the IEEE International Conference on Computer Vision*. 2019.


GANs for 3D Objects

Chair




3D-GAN


Car



Sofa



Table



Generative
Design

Wu, Jiajun, et al. "Learning a probabilistic latent space of object shapes via 3d generative-adversarial modeling." *Advances in neural information processing systems*. 2016.

Companies Using GANs



Next-gen
Photoshop



Text
Generation



Data
Augmentation





Image Filters



Super-resolution

Summary

- GANs' performance is rapidly improving
- Huge opportunity to work in this space!
- Major companies are using them





deeplearning.ai

Intuition Behind GANs

Outline

- The goal of the generator and the discriminator
- The competition between them




Generative Adversarial Network

Generator learns to make *fakes*
that look *real*



Discriminator learns to distinguish
real from *fake*



Generative Adversarial Network

Discriminator learns to distinguish
real from *fake*



Generative Adversarial Network

Generator learns to make *fakes*
that look **real**

Discriminator learns to distinguish
real from *fake*



Generative Adversarial Network

Generator learns to make *fakes*
that look *real*



Doesn't know how
it should look



Generative Adversarial Network

Discriminator learns to distinguish
real from *fake*



The Game Is On!



5% Real




40% Real



80% Real




The Game is On!



The Game Is On!



30% Real



60% Real




95% Real

The Game Is On!




30% Real



60% Real




Wrong!



95% Real

Won't fool me
again

The Game Is On!



Summary

- The **generator's** goal is to fool the discriminator
- The **discriminator's** goal is to distinguish between real and fake
- They learn from the competition with each other
- At the end, *fakes* look **real**






deeplearning.ai

Discriminator

Outline

- Review of classifiers
- The role of classifiers in terms of probability
- Discriminator



Classifiers

Distinguish between different classes



Turtle

Bird


Cat

Dog


Fish

Classifiers


Distinguish between different classes



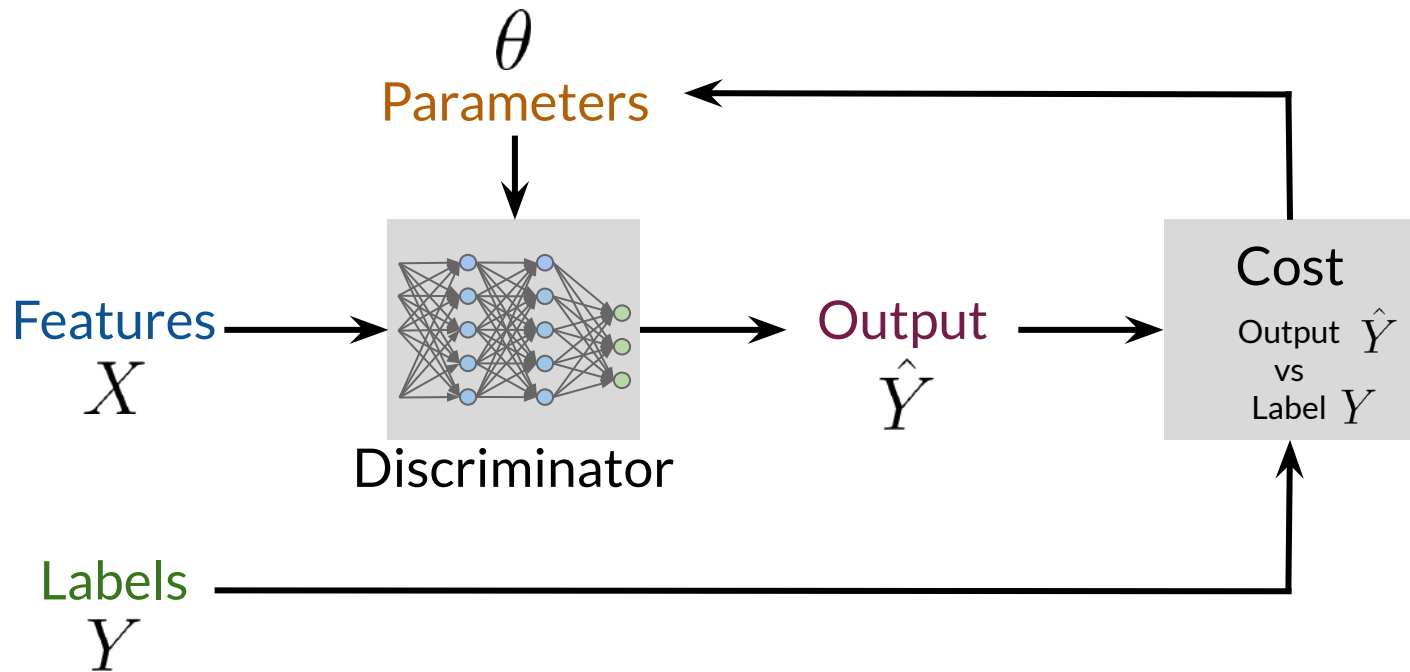
Neural Networks



Neural Networks




Classifiers (training)




Classifiers



$$P\left(\begin{array}{c} \text{Turtle} \\ \text{Bird} \\ \text{Cat} \\ \text{Dog} \\ \text{Fish} \end{array} \mid \text{Image} \right)$$





Classifiers



Discriminator



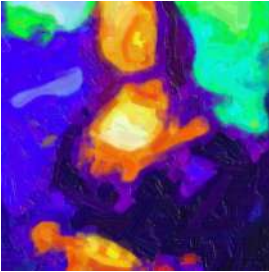
Discriminator



Discriminator

$$P\left(\begin{array}{c} \text{Fake} \\ \text{Class} \end{array} \mid \begin{array}{c} X \\ \text{Features} \end{array} \right)$$

Discriminator

$$P\left(\begin{array}{c} \text{Fake} \\ \text{Class} \end{array} \mid \begin{array}{c} \text{Features} \end{array}\right) = 0.85 \rightarrow \boxed{\text{Fake}}$$


Summary

- The **discriminator** is a classifier
- It learns the probability of class Y (**real** or *fake*) given features X
- The probabilities are the feedback for the **generator**





deeplearning.ai

Generator

Outline

- What the generator does
- How it improves its performance
- Generator in terms of probability



Generator

Turtle

Generates examples of the class

Bird


Cat

Dog


Fish



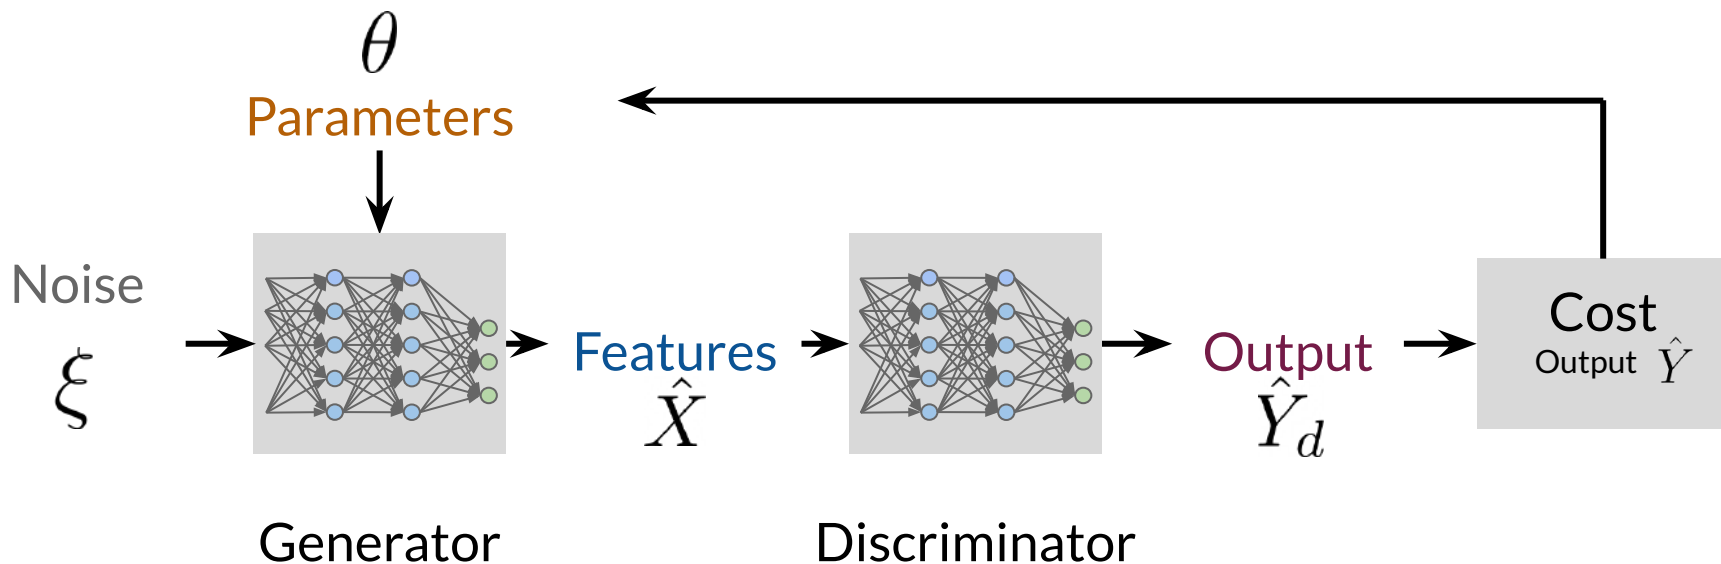
Neural Networks




Neural Networks



Generator: Learning



Sampling



Generator

$$P(\text{Image} \mid \text{Class})$$

Turtle


Bird

Cat

Dog

Fish

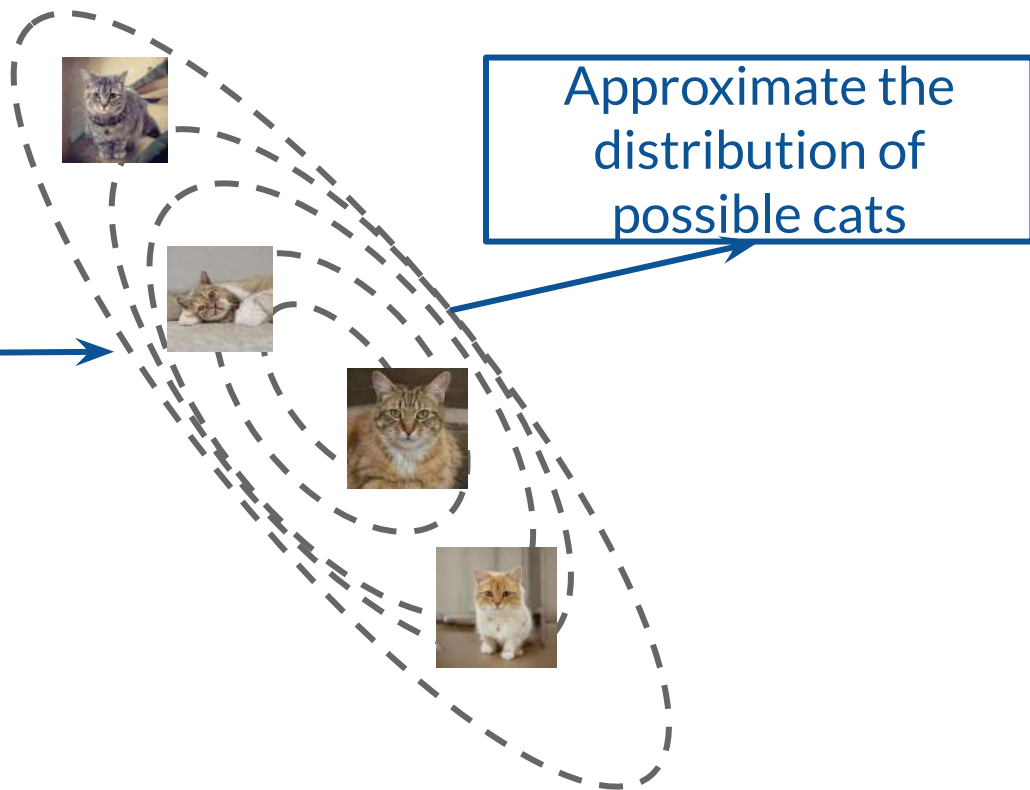
Generator



Generator

$P(X)$

Features



Images available from: <http://thesecatsdonotexist.com/>

Summary

- The **generator** produces fake data
- It learns the probability of features X
- The **generator** takes as input noise (random features)





deeplearning.ai

BCE Cost Function

Outline

- Binary Cross Entropy (BCE) Loss equation by parts
- How it looks graphically



BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

Prediction

Label

Features

Parameters

Average loss of the whole batch

The diagram illustrates the components of the Binary Cross-Entropy (BCE) cost function. The formula is shown as $J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$. Annotations include: a purple circle around the summation term with an arrow pointing to a box labeled 'Average loss of the whole batch'; a grey arrow from $h(x^{(i)}, \theta)$ pointing to 'Prediction'; a grey arrow from $y^{(i)}$ pointing to 'Label'; a grey arrow from $x^{(i)}$ pointing to 'Features'; and a grey arrow from θ pointing to 'Parameters'.

BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

$y^{(i)}$	$h(x^{(i)}, \theta)$	$y^{(i)} \log h(x^{(i)}, \theta)$
0	any	0
1	0.99	~ 0
1	~ 0	-inf

Relevant when
the label is 1

BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

$y^{(i)}$	$h(x^{(i)}, \theta)$	$(1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))$
1	any	0
0	0.01	~0
0	~1	-inf

Relevant when
the label is 0


BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

Ensures that the cost is always greater or equal to 0



BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$



BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$



Summary

- The BCE cost function has two parts (one relevant for each class)
- Close to zero when the label and the prediction are similar
- Approaches infinity when the label and the prediction are different






deeplearning.ai


Putting It All Together

Outline

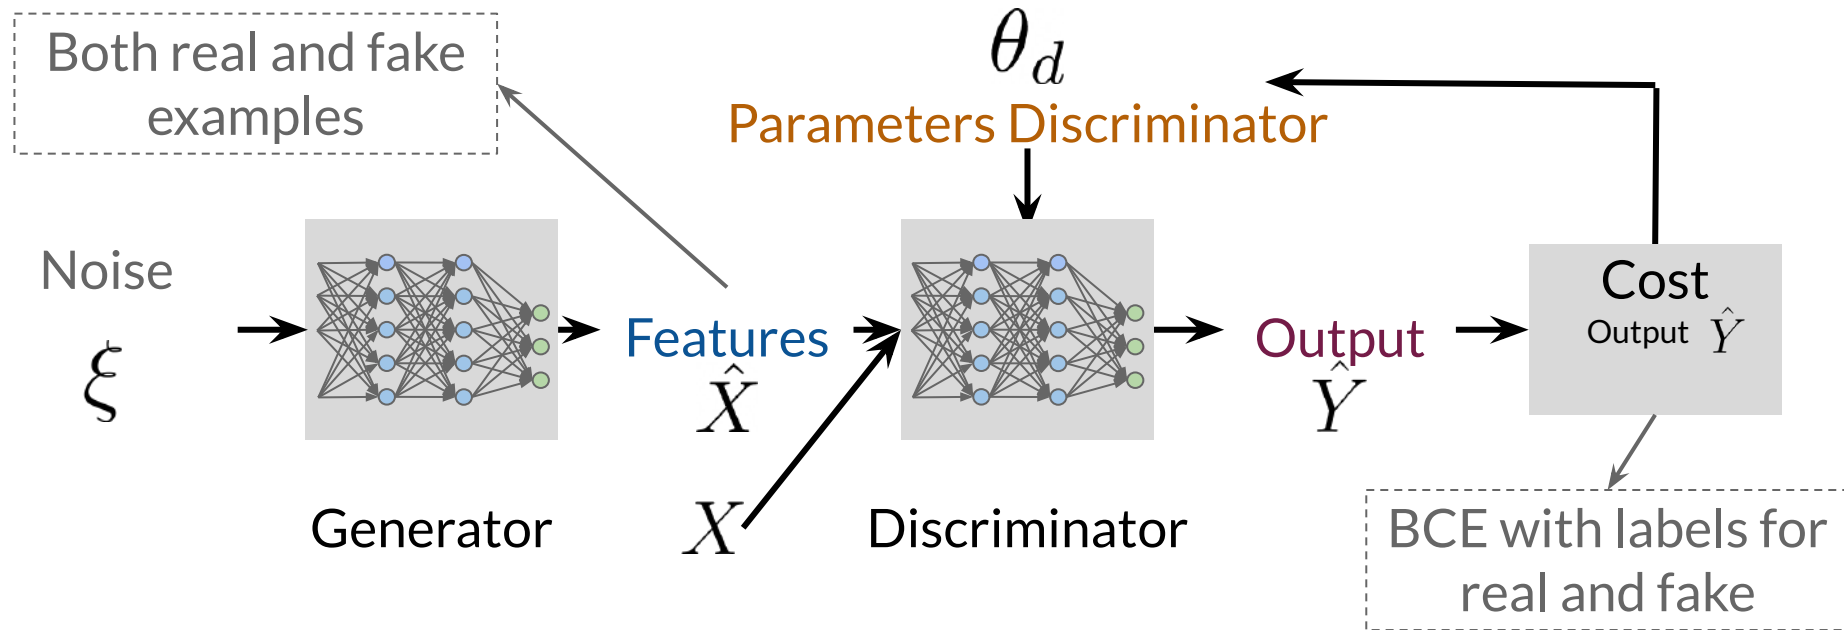
- How the whole architecture looks
- How to train GANs



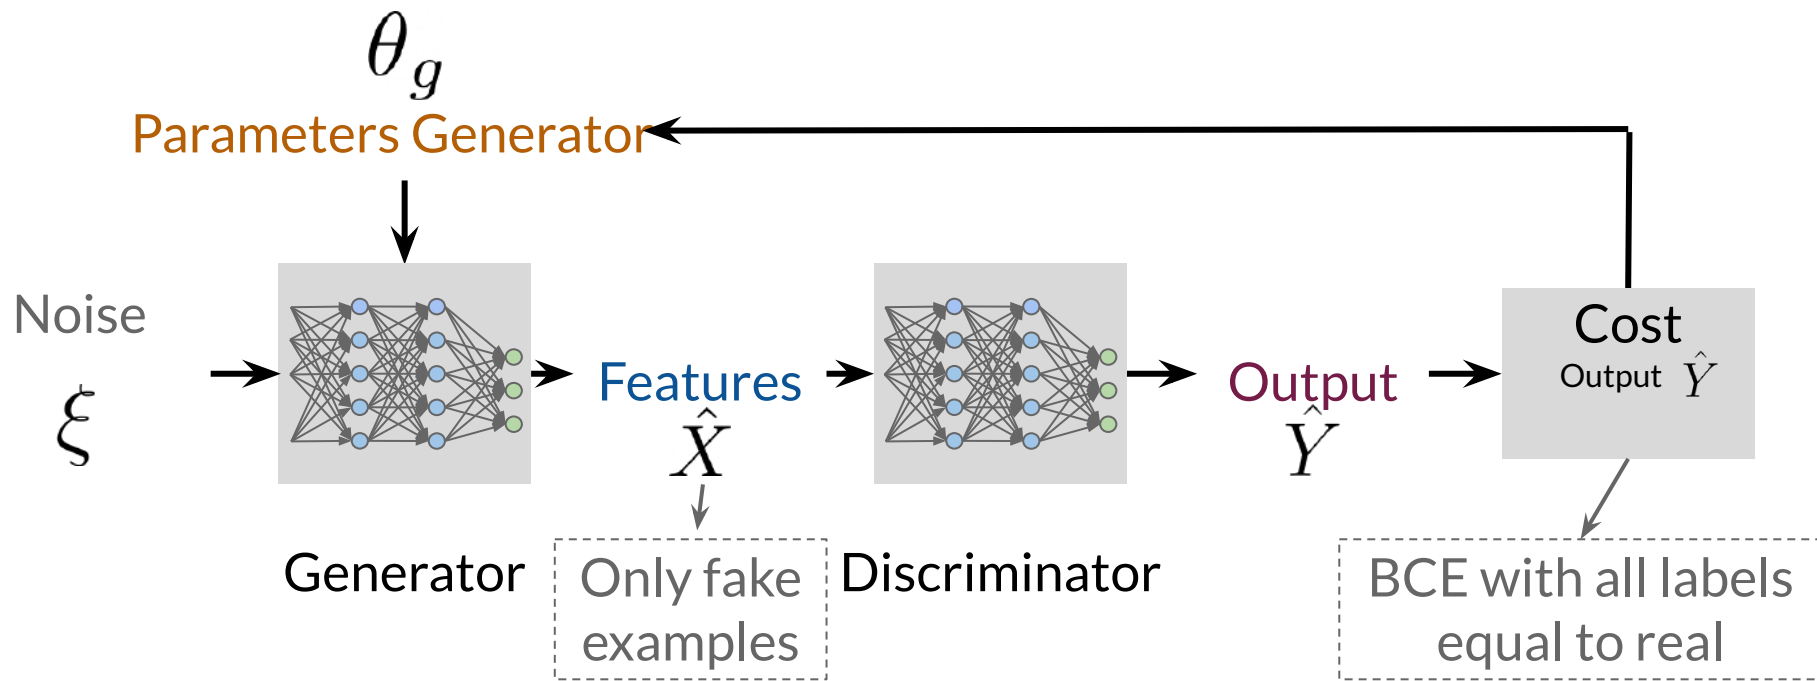
GANs Model



Training GANs: Discriminator



Training GANs: Generator




Training GANs



Summary

- GANs train in an alternating fashion
- The two models should always be at a similar “skill” level





deeplearning.ai

Intro to PyTorch (Optional)

Outline

- Comparison with TensorFlow
- Defining Models
- Training



PyTorch vs TensorFlow

PyTorch

Imperative, computations on the go

```
A, B = 1, 2
C = A + B
print(C)
```

3

Dynamic Computational Graphs

TensorFlow

Symbolic, first define and then compile

```
C = A + B
f = compile(C)
print(f(A = 1, B = 2))
```

3

Static Computational Graphs

Tensorflow > 2.0 moves toward PyTorch by including **Eager Execution**

PyTorch vs TensorFlow

PyTorch

TensorFlow

Currently very similar frameworks!

Tensorflow > 2.0 moves toward PyTorch by
including Eager Execution

Defining Models in PyTorch

```
import torch
from torch import nn
```

→ Custom layers for DL

```
class LogisticRegression(nn.Module):
    def __init__(self, in):
        super().__init__()
        self.log_reg = nn.Sequential(
            nn.Linear(in, 1),
            nn.Sigmoid()
        )
    def forward(self, x):
        return self.log_reg(x)
```

Define the model as a class

Initialization method with parameters

Definition of the architecture

Forward computation of the model
with inputs x

Training Models In PyTorch

```
model = LogisticRegression(16)
```

Initialization of the model

```
criterion = nn.BCELoss()
```

Cost function

```
optimizer = torch.optim.SGD(model.parameters(), lr=0.01)
```

Optimizer

```
for t in range(n_epochs):
```

Training loop for number of epochs

```
    y_pred = model(x)
```

```
    loss = criterion(y_pred, y)
```

Forward propagation

```
    optimizer.zero_grad()
```

```
    loss.backward()
```

```
    optimizer.step()
```

Optimization step

Summary

- PyTorch makes computations on the run
- Dynamic computational graphs in Pytorch
- Just another framework, and similar to Tensorflow!

