




deeplearning.ai

Conditional Generation: Intuition

Outline

- Unconditional generation
- Conditional vs. unconditional generation



Unconditional Generation

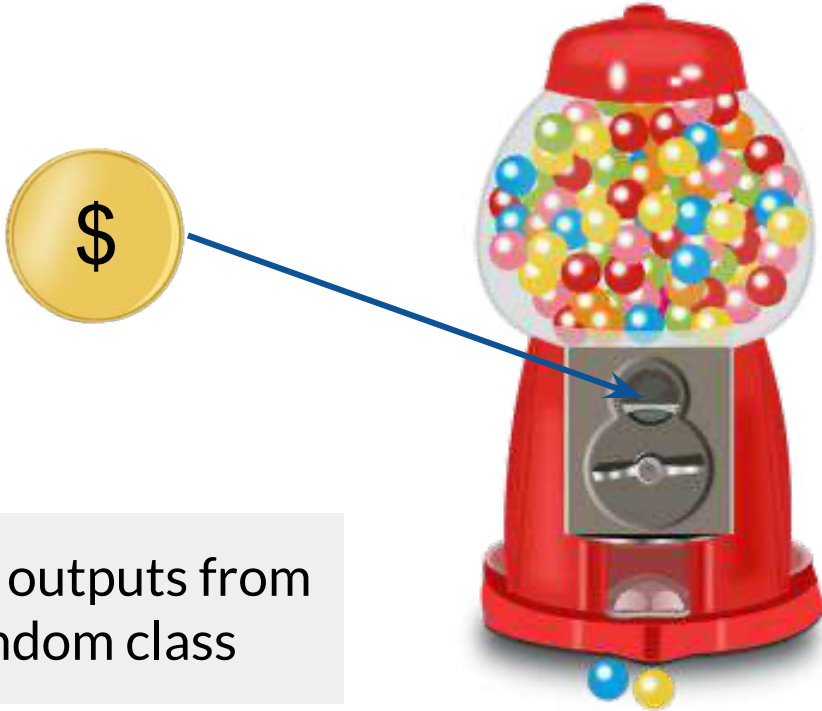
You get outputs from
a random class

Unconditional Generation

You get outputs from
a random class




Unconditional Generation

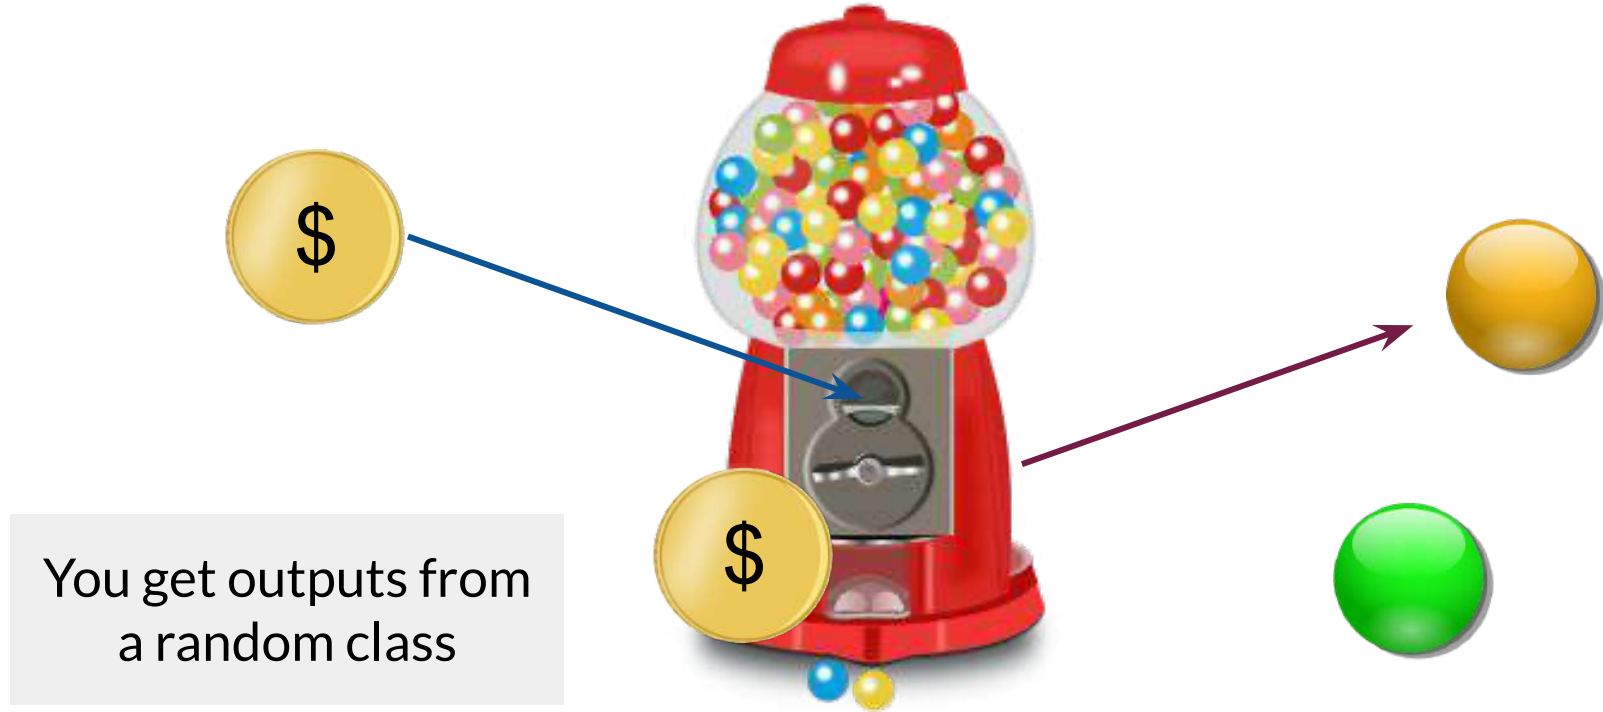


You get outputs from
a random class


Unconditional Generation



Unconditional Generation



Unconditional Generation

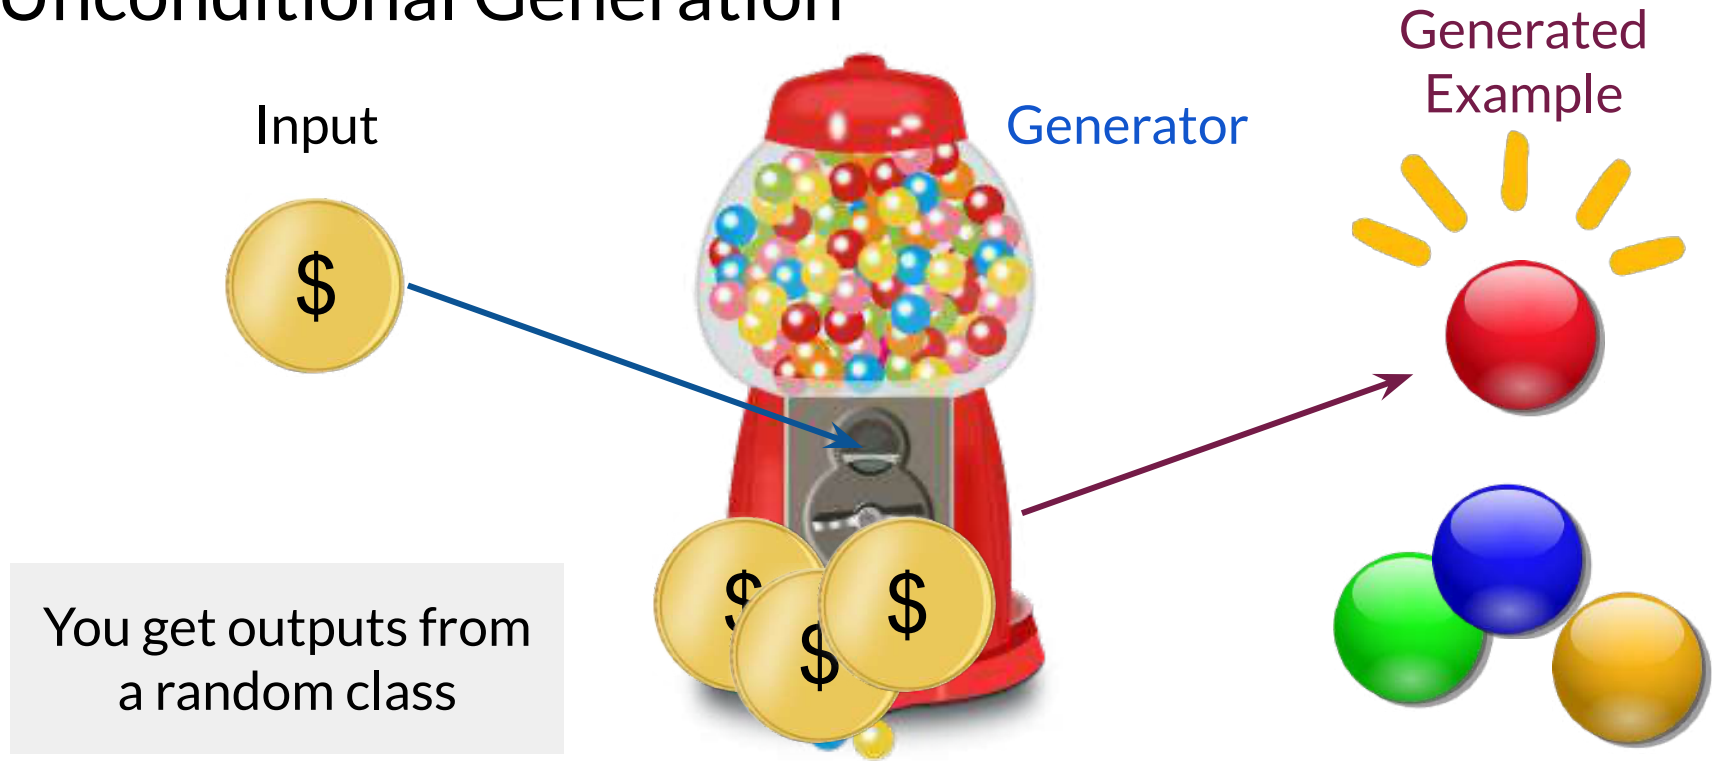


Unconditional Generation



You get outputs from
a random class


Unconditional Generation



Conditional Generation

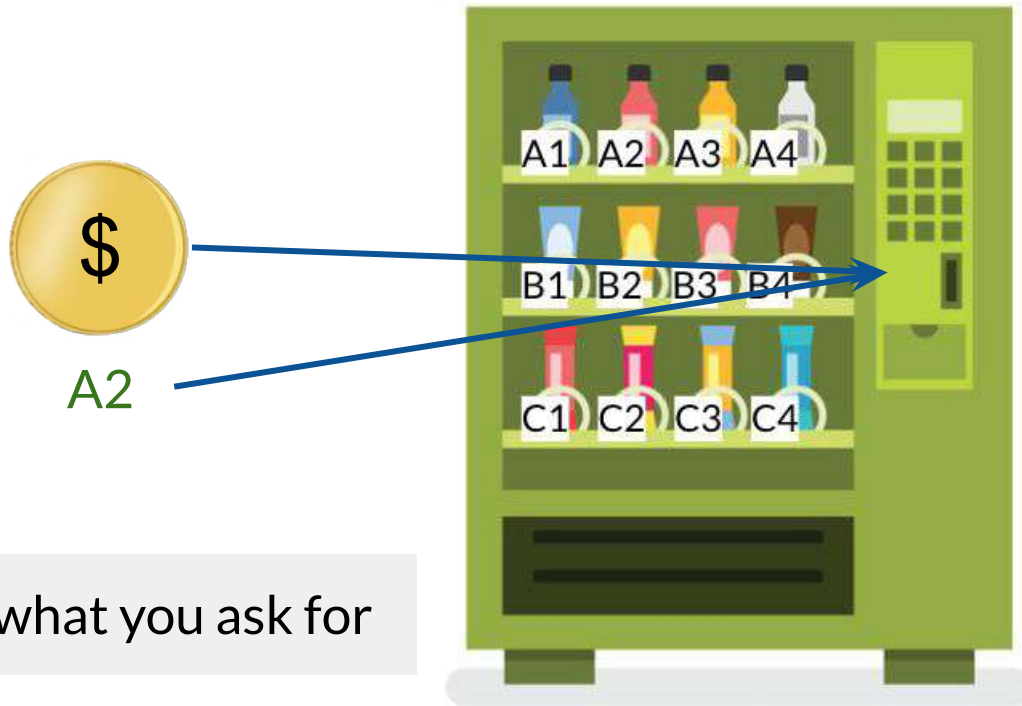
You get what you ask for

Conditional Generation




You get what you ask for

Conditional Generation



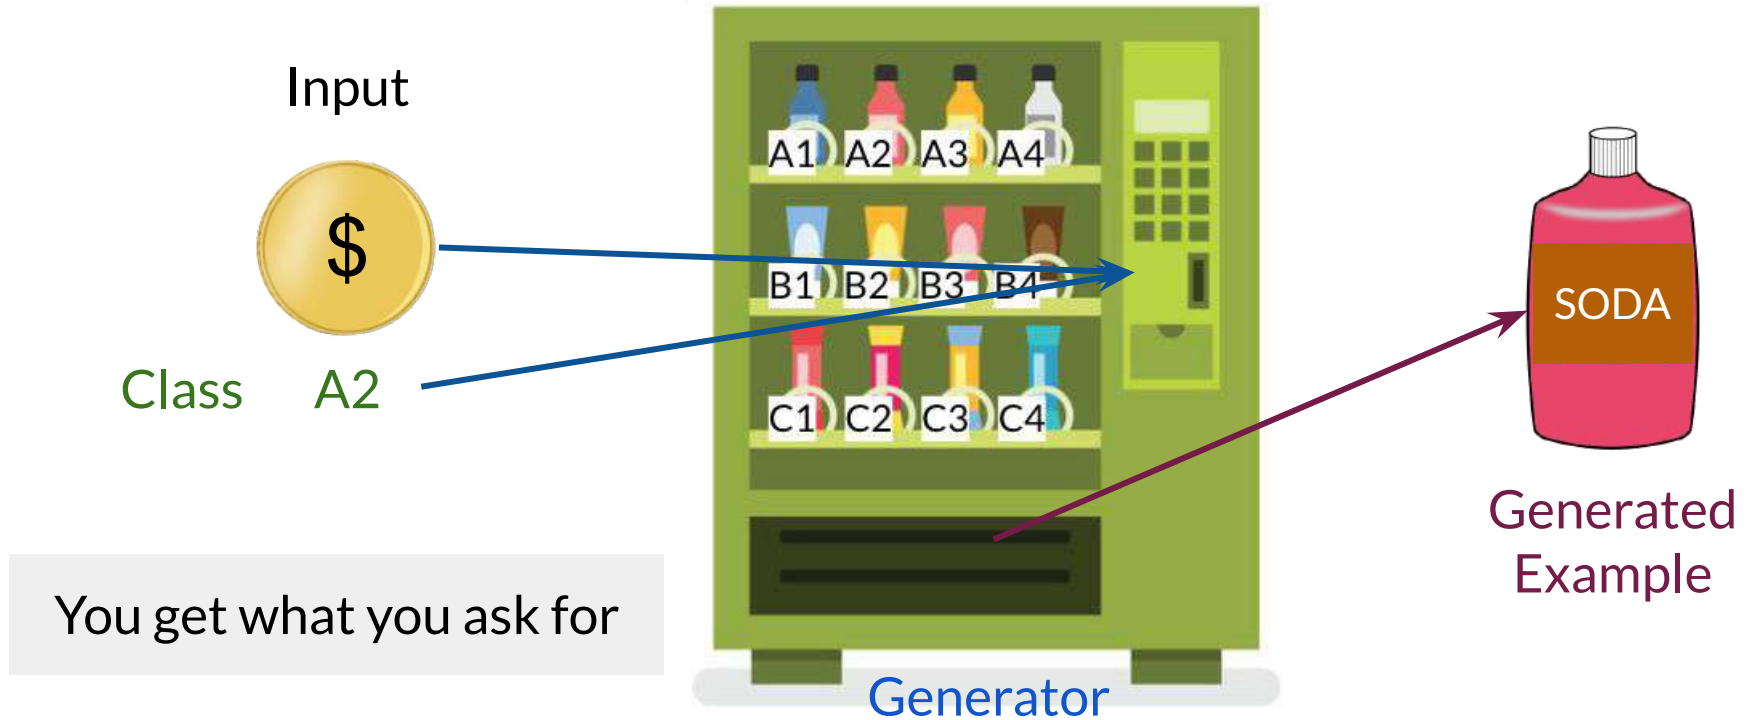
You get what you ask for

Conditional Generation



You get what you ask for

Conditional Generation



Conditional vs. Unconditional Generation



Conditional vs. Unconditional Generation

Conditional

Unconditional

Examples from **the classes
you want**


Examples from *random classes*

Conditional vs. Unconditional Generation

Conditional	Unconditional
Examples from the classes you want	Examples from <i>random classes</i>
Training dataset needs to be labeled	Training dataset <i>doesn't need to be labeled</i>

Summary

- Conditional generation requires labeled datasets
- Examples can be generated for the selected class






deeplearning.ai


Conditional Generation: Inputs

Outline

- How to tell the generator what type of example to produce
- Input representation for the discriminator




Generator Input




Generator Input

Noise vector




Class (one-hot) vector




Generator Input

Noise vector




Class (one-hot) vector



Husky


Generator Input

Noise vector



Randomness in the
generation


Class (one-hot) vector



Husky


Generator Input

Noise vector



Randomness in the generation


Class (one-hot) vector




Husky

Control in the generation

Generator Input




Generator Input



Generator Input

Output




Discriminator Input




Discriminator


Discriminator Input



Discriminator Input



Discriminator Input



Discriminator Input



Golden Retriever




Output

98%
Real

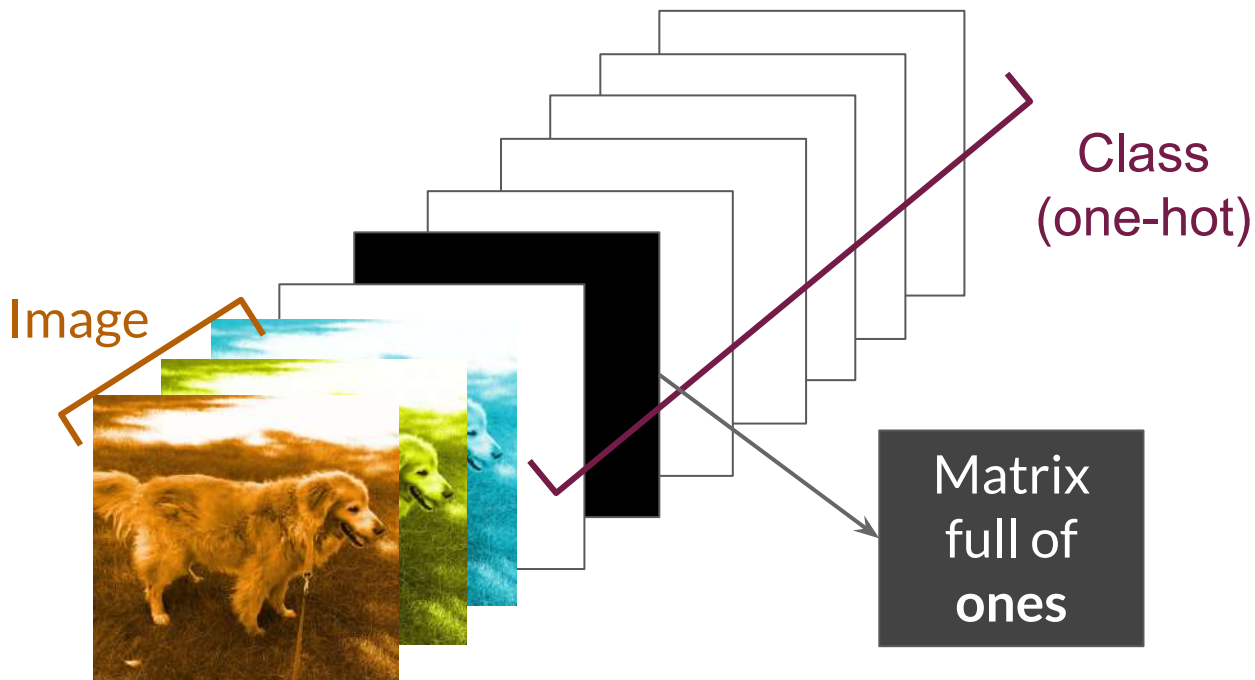
Discriminator Input




Discriminator Input



Discriminator Input



Discriminator Input



Summary

- The class is passed to the generator as one-hot vectors
- The class is passed to the discriminator as one-hot matrices
- The size of the vector and the number of matrices represent the number of classes






deeplearning.ai

Controllable Generation

Outline

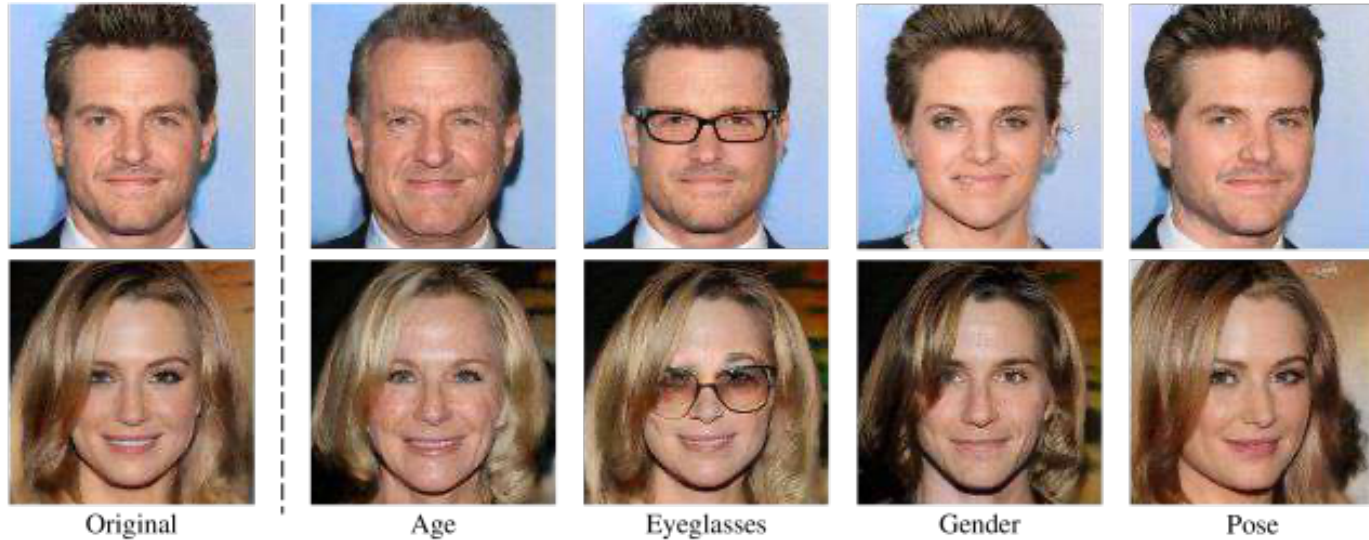
- What is controllable generation
- How it compares to conditional generation



Controllable Generation

Change specific features of the output

Controllable Generation



Change specific features of the output

Available from: <https://arxiv.org/abs/1907.10786>

Controllable Generation

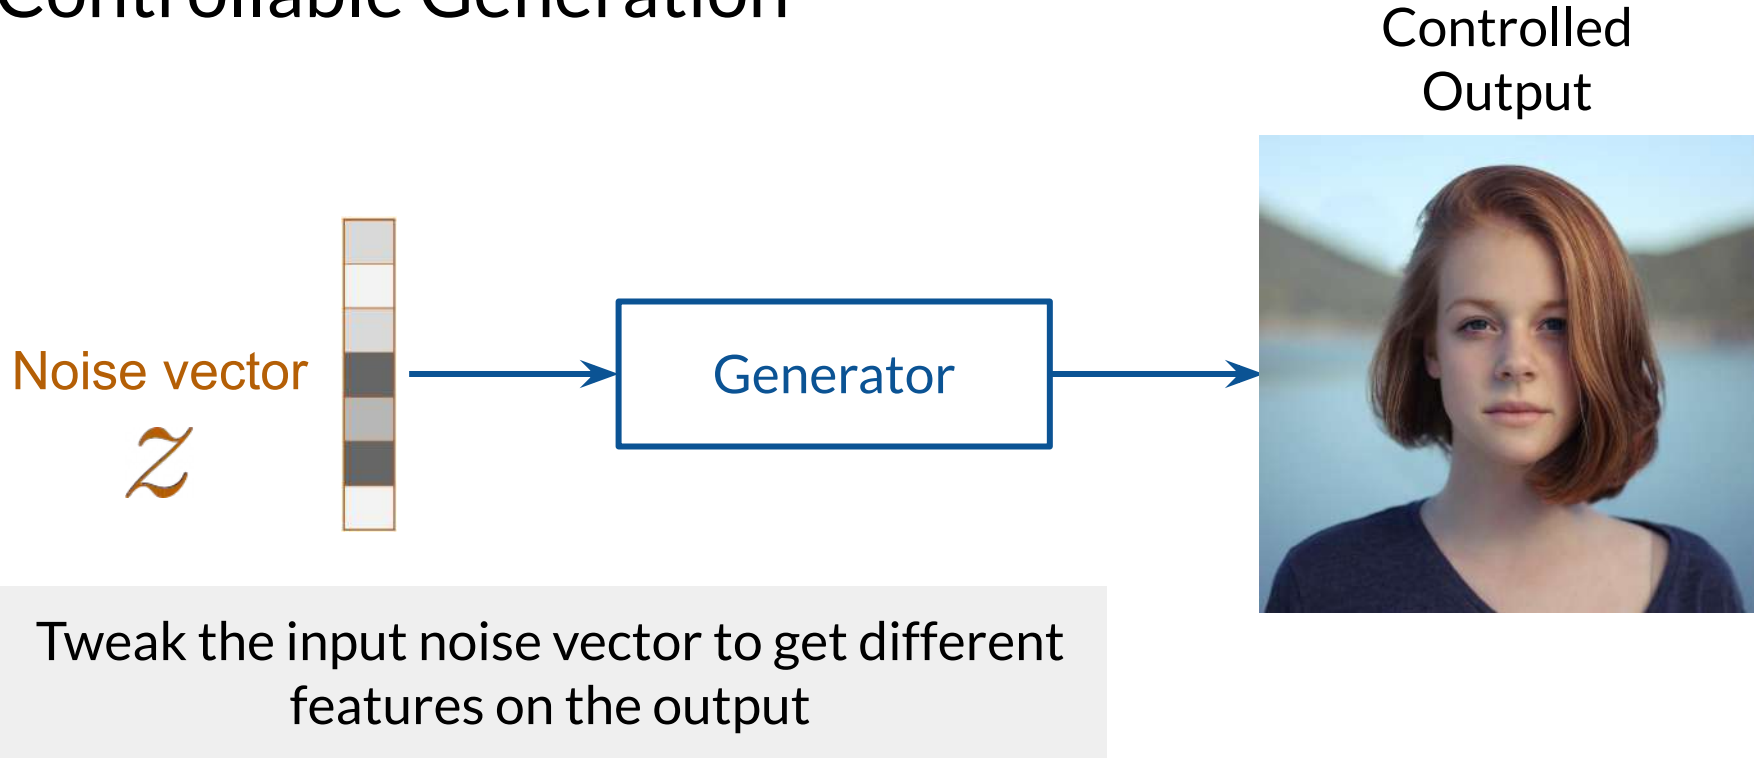
Noise vector

z

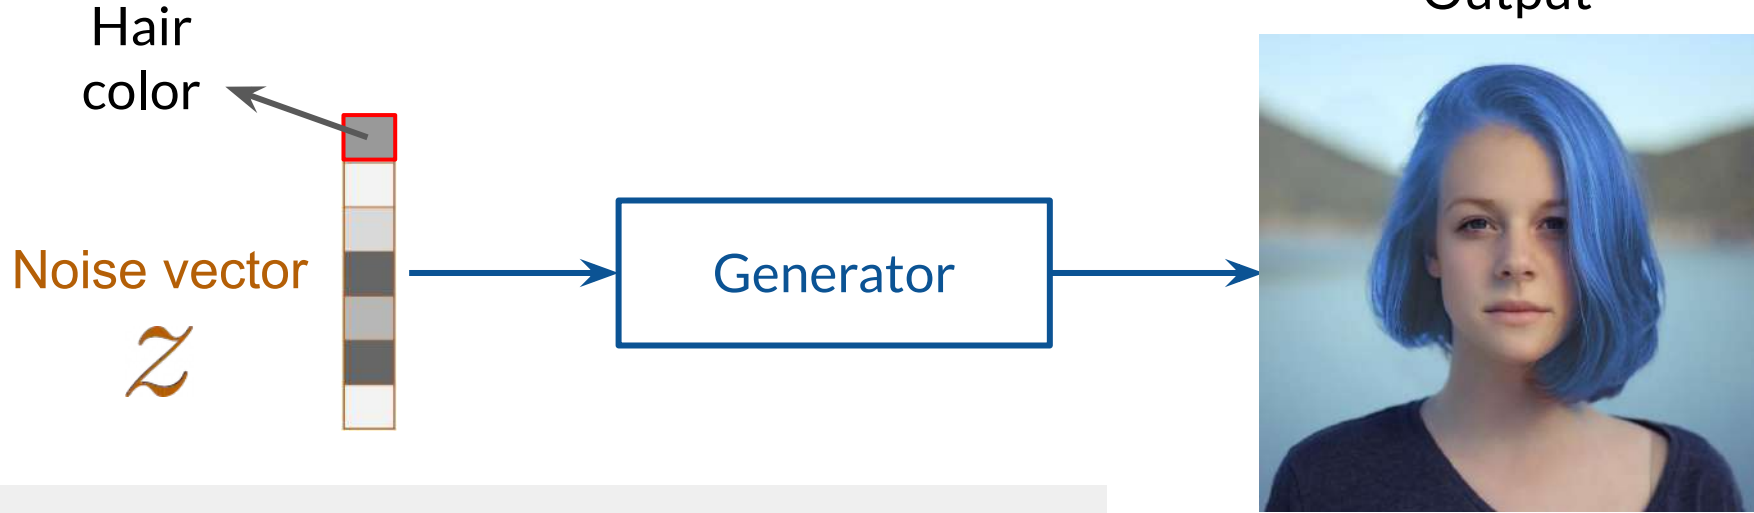


Tweak the input noise vector to get different features on the output

Controllable Generation



Controllable Generation



Tweak the input noise vector to get different features on the output

Controllable Generation vs. Conditional Generation

Controllable

Conditional

Controllable Generation vs. Conditional Generation

Controllable

Conditional

Examples with the **features**
that you want

Examples from *the classes you*
want

Controllable Generation vs. Conditional Generation

Controllable

Examples with the **features**
that you want

Training dataset **doesn't need**
to be labeled

Conditional

Examples from *the classes you*
want

Training dataset *needs to be*
labeled

Controllable Generation vs. Conditional Generation

Controllable

Examples with the **features that you want**

Training dataset **doesn't need to be labeled**

Manipulate the z vector
input

Conditional


Examples from *the classes you want*

Training dataset *needs to be labeled*

Append a class vector to the
input

Summary

- Controllable generation lets you control the features of the generated outputs
- It does not need a labeled training dataset
- The input vector is tweaked to get different features on the output






deeplearning.ai



Vector Algebra in the Z-Space

Outline

- Interpolation in the Z-space
- Modifying the noise vector z to control desired features




Interpolation Using the Z-Space




How an image morphs into another

Interpolation Using the Z-Space



Interpolation Using the Z-Space




$g(v_1)$



$g(v_2)$



Interpolation Using the Z-Space




$g(v_1)$



$g(v_2)$



Interpolation Using the Z-Space




Intermediate images using
the Z-space


$g(v_1)$ \longleftrightarrow $g(v_2)$



Z-Space and Controllable Generation




Z-Space and Controllable Generation




Output

Z-Space and Controllable Generation




Output


Z-Space and Controllable Generation



Z-Space and Controllable Generation



Z-Space and Controllable Generation




Original output

$$g(v_1) \longrightarrow$$



Z-Space and Controllable Generation




Original output

$$g(v_1) \longrightarrow$$



Z-Space and Controllable Generation



Original output

$$g(v_1) \rightarrow$$




Controlled output

$$g(v_1 + d) \rightarrow$$



Summary

- To control output features, you need to find directions in the Z-space
- To modify your output, you move around in the Z-space





deeplearning.ai

Challenges with Controllable Generation


Outline

- Output feature correlation
- Z-space entanglement



Feature Correlation

Uncorrelated
Features




Add beard



Feature Correlation


Uncorrelated
Features



Add beard




Correlated
Features



Add beard

Make more
masculine



Z-Space Entanglement

Noise vector

z



Glasses

Beard


Hair

Eyes


Age

Output
Features


Z-Space Entanglement




Z-Space Entanglement




Z-Space Entanglement



Z-Space Entanglement




Z-Space Entanglement



It is not possible to control single output features

Z-Space Entanglement



It is not possible to control single output features

Summary

- When trying to control one feature, others that are correlated change
- Z-space entanglement makes controllability difficult, if not impossible
- Entanglement happens when z does not have enough dimensions






deeplearning.ai

Classifier Gradients

Outline

- How to use classifiers to find directions in the Z-space
- Requirements to use this method




Classifier Gradients

Classifier Gradients

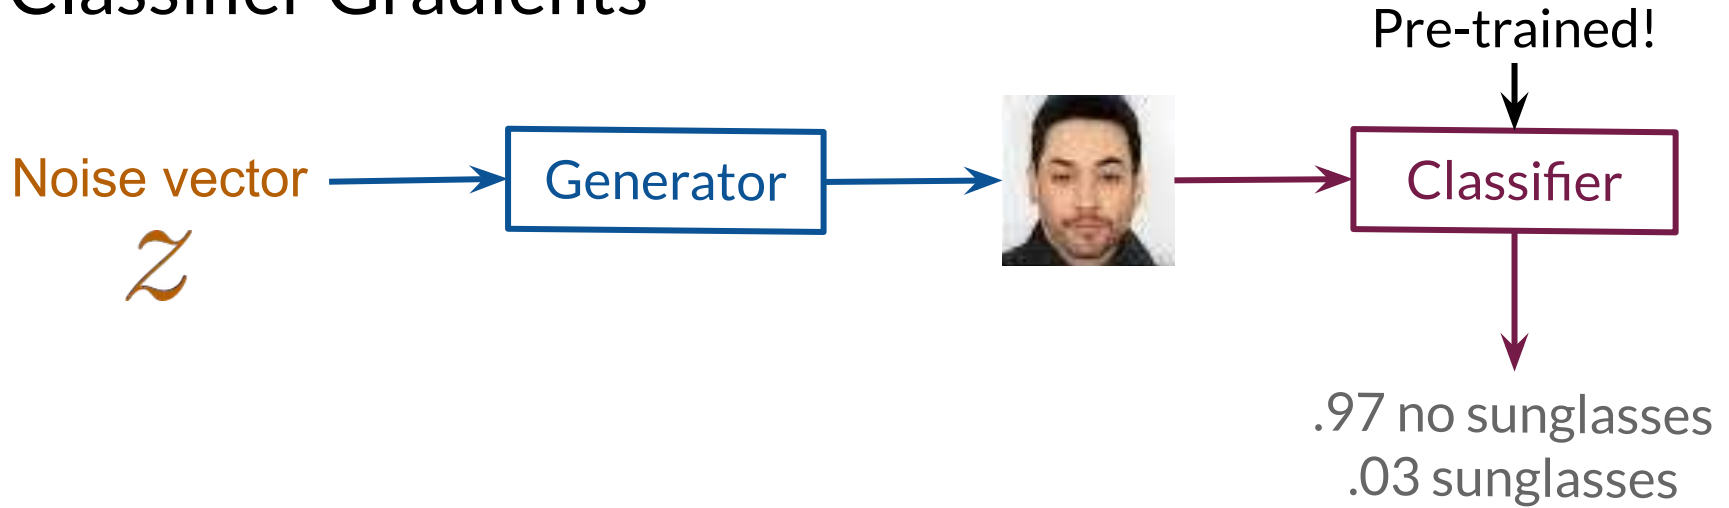
Noise vector

z

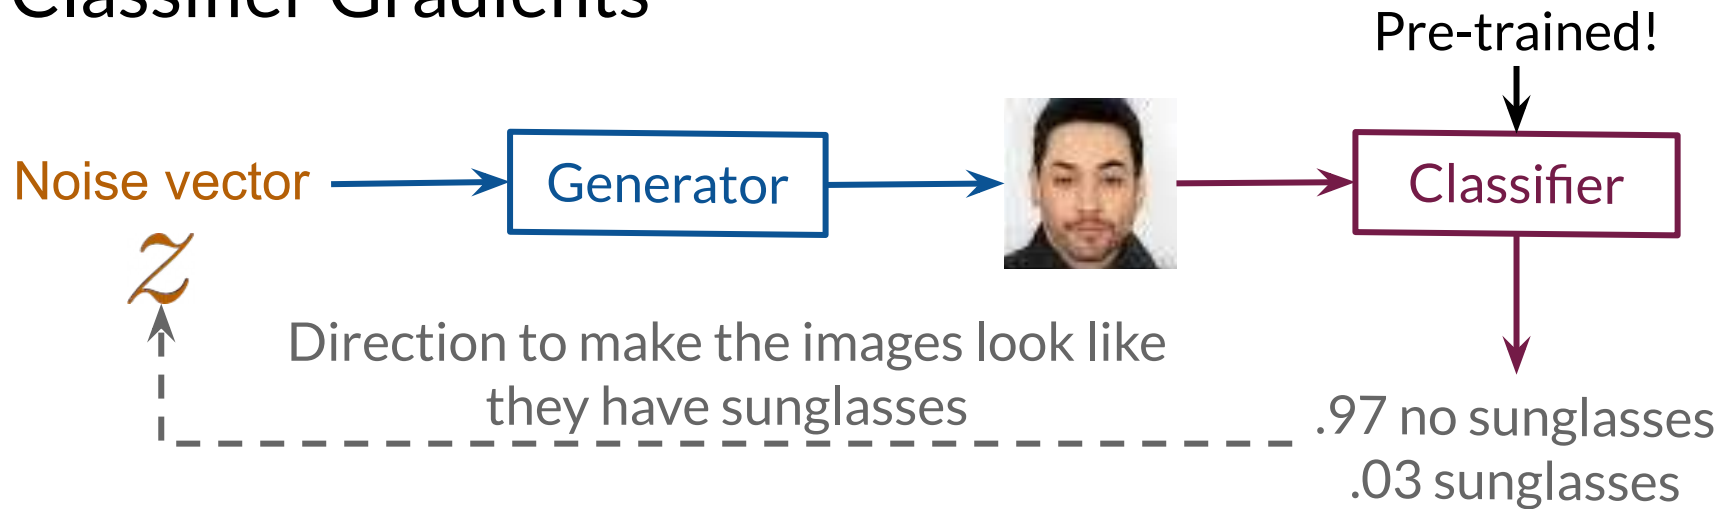
Classifier Gradients



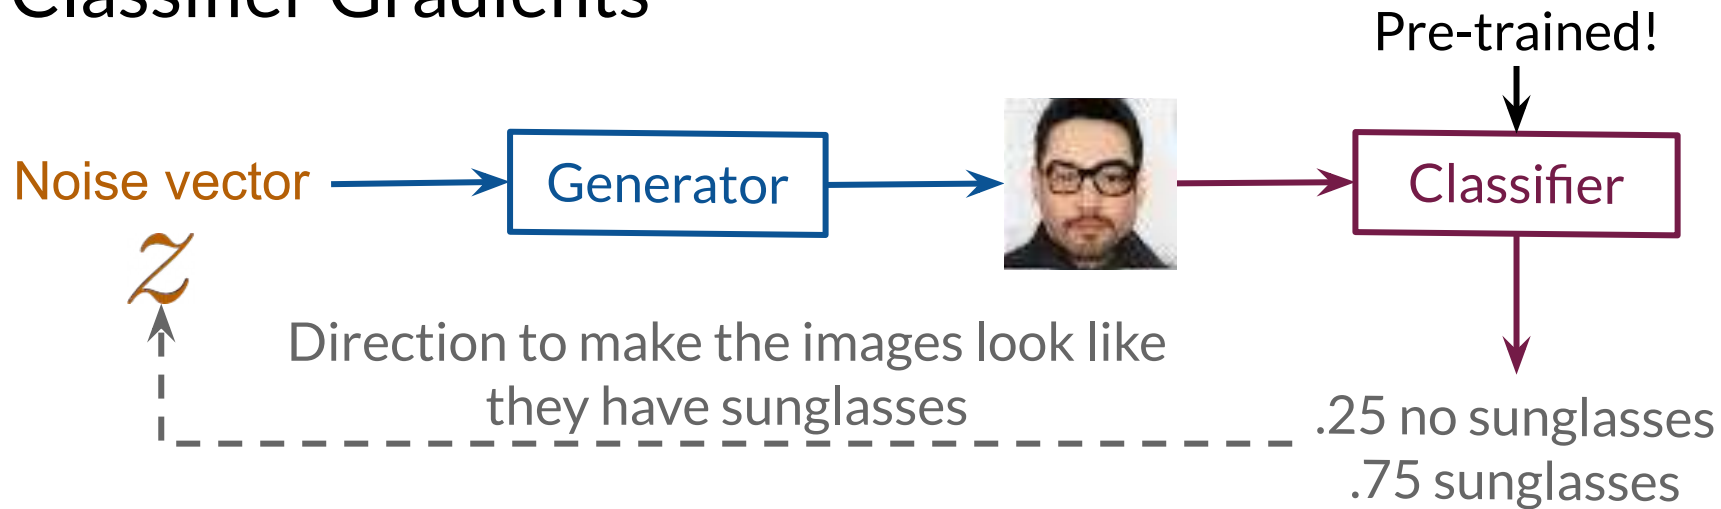
Classifier Gradients



Classifier Gradients

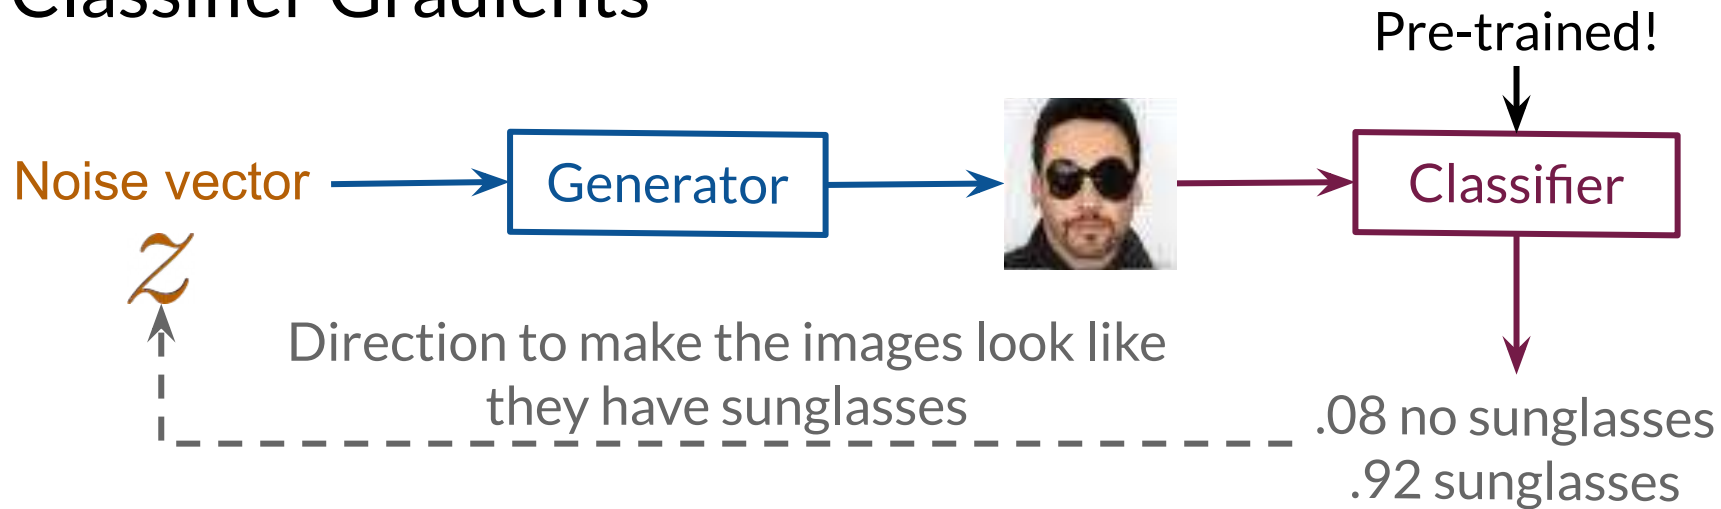


Classifier Gradients



Modify **just** the **noise vector** until the feature emerges

Classifier Gradients



Modify **just** the **noise vector** until the feature emerges

Summary

- Classifiers can be used to find directions in the Z-space
- To find directions, the updates are done just to the noise vector






deeplearning.ai

Disentanglement

Outline


- What a disentangled Z-space means
- Ways to encourage disentangled Z-spaces



Disentangled Z-Space

$$v_1 = [1, 2, 3, \dots]$$

$$v_2 = [5, 6, 7, \dots]$$




Disentangled Z-Space

$$v_1 = [\overset{z_1}{1}, 2, 3, \dots]$$


$$v_2 = [5, 6, 7, \dots]$$

Hair
color



Disentangled Z-Space


$$\begin{array}{c} z_1 \quad z_2 \\ v_1 = [\text{1}, \text{2}, 3, \dots] \\ v_2 = [\text{5}, \text{6}, 7, \dots] \\ \text{Hair color} \quad \text{Hair length} \end{array}$$



Disentangled Z-Space

$$\begin{array}{c} z_1 \quad z_2 \\ v_1 = [\text{1}, \text{2}, 3, \dots] \\ v_2 = [\text{5}, \text{6}, 7, \dots] \end{array}$$


Hair color Hair length



Disentangled Z-Space

$$\begin{array}{c} z_1 \quad z_2 \\ v_1 = [\text{1}, \text{2}, 3, \dots] \\ v_2 = [\text{5}, \text{6}, 7, \dots] \end{array}$$

Hair color Hair length




Disentangled Z-Space


$$\begin{array}{c} z_1 \quad z_2 \\ v_1 = [\text{1}, \text{2}, 3, \dots] \\ v_2 = [\text{5}, \text{6}, 7, \dots] \end{array}$$

Hair color Hair length


Latent factors of variation



Disentangled Z-Space




Disentangled Z-Space




Changes to one feature
don't affect the others

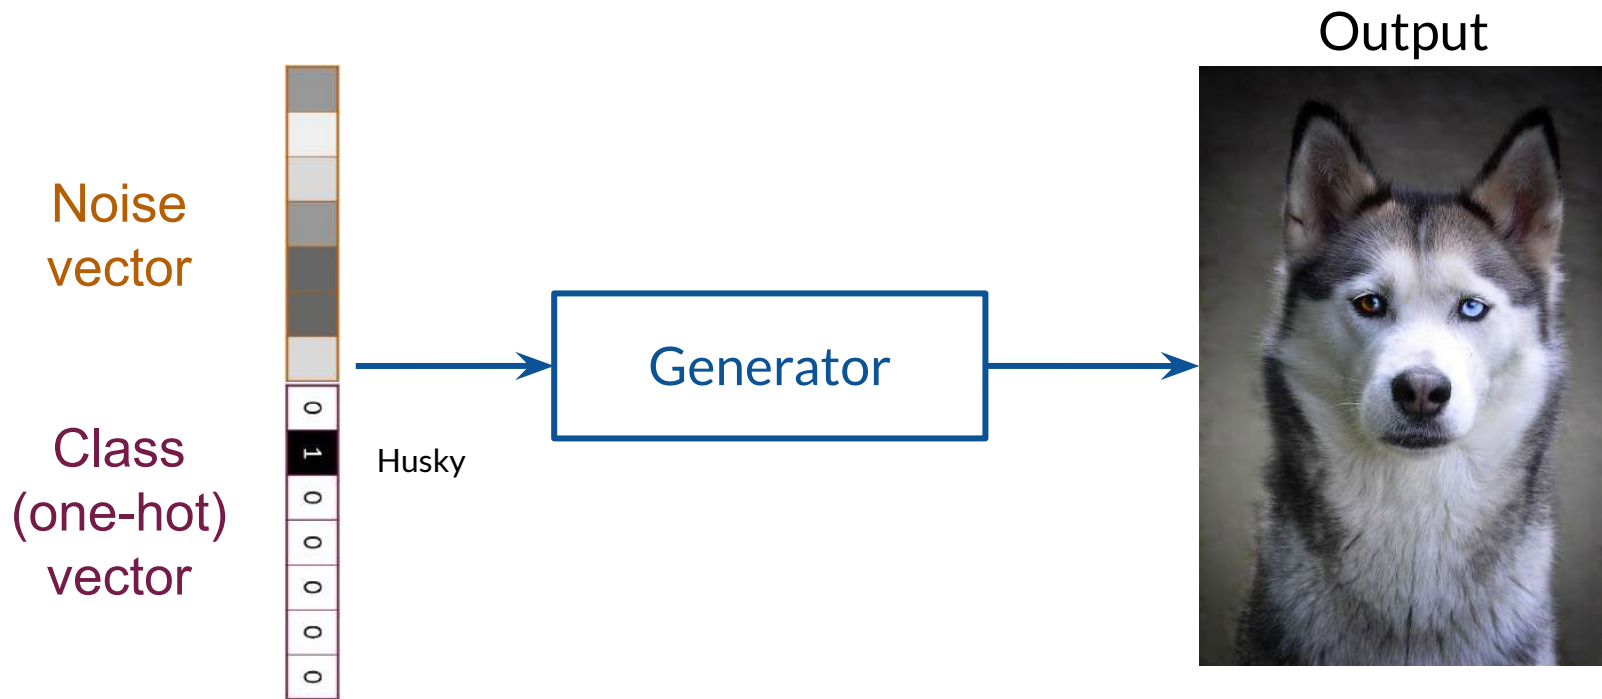
Disentangled Z-Space



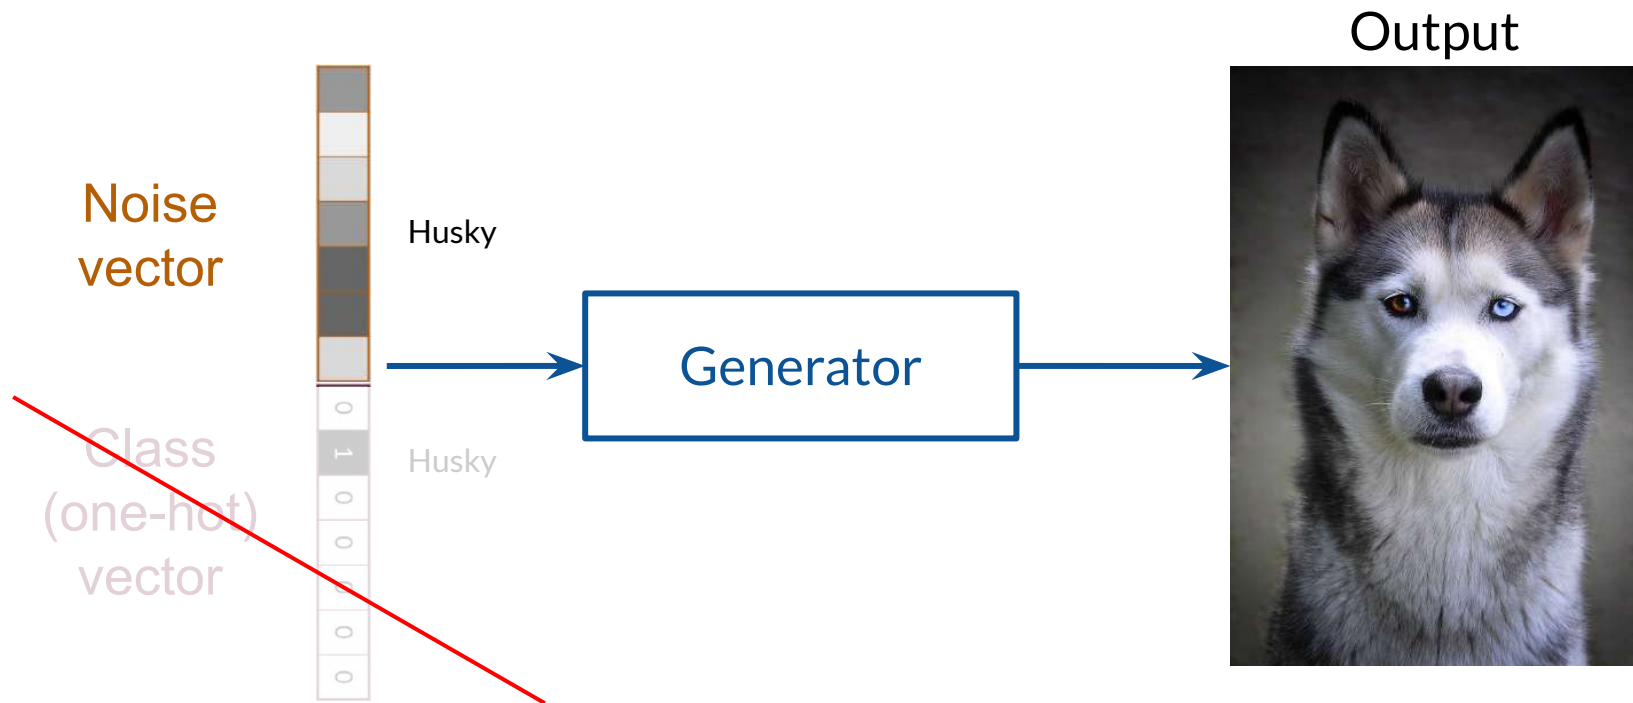
Changes to one feature
don't affect the others



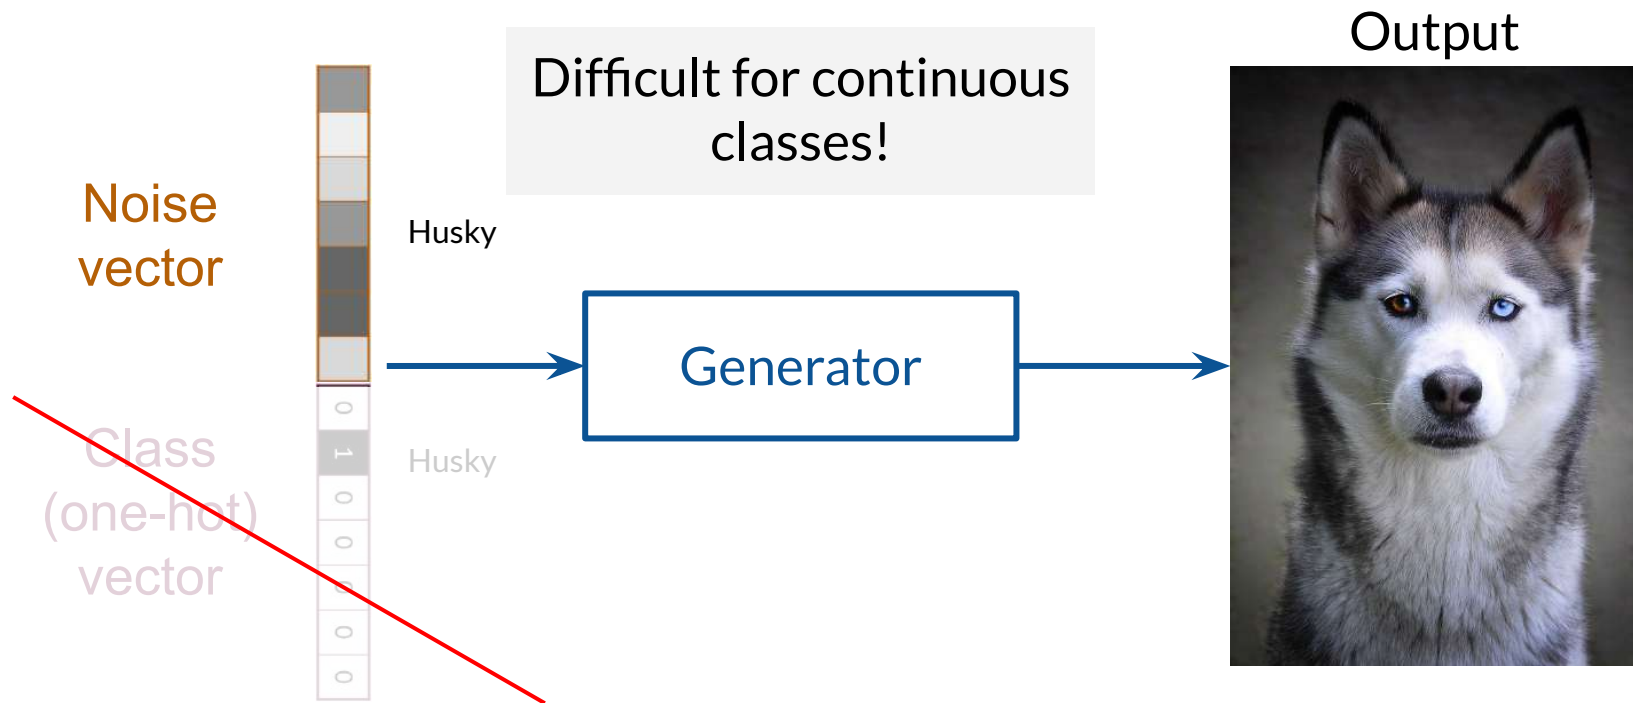
Encourage Disentanglement: Supervision



Encourage Disentanglement: Supervision



Encourage Disentanglement: Supervision



Encourage Disentanglement: Loss Function

$$v_1 = [1, 2, 3, \dots]$$

$$v_2 = [5, 6, 7, \dots]$$

Encourage Disentanglement: Loss Function

$$v_1 = [1, 2, 3, \dots]$$

$$v_2 = [5, 6, 7, \dots]$$

$$L_{\text{new}} = \boxed{L_{\text{original}}} + \boxed{\text{reg}_d}$$

Original loss Regularization

Encourage Disentanglement: Loss Function

$$v_1 = [1, 2, 3, \dots]$$

$$v_2 = [5, 6, 7, \dots]$$

$$L_{\text{new}} = \boxed{L_{\text{original}}} + \boxed{\text{reg}_d}$$

Original loss Regularization

Can be any loss function
(e.g. BCE, W-Loss)

Encourage Disentanglement: Loss Function


$$\begin{array}{cc} & z_1 & z_2 \\ v_1 = [& 1, & 2, & 3, \dots] \\ v_2 = [& 5, & 6, & 7, \dots] \end{array}$$

$$L_{\text{new}} = \boxed{L_{\text{original}}} + \boxed{\text{reg}_d}$$

Original loss Regularization

Can be any loss function
(e.g. BCE, W-Loss)

Encourage Disentanglement: Loss Function



$$L_{\text{new}} = \boxed{L_{\text{original}}} + \boxed{\text{reg}_d}$$

Original loss Regularization

Can be any loss function
(e.g. BCE, W-Loss)

Summary

- Disentangled Z-spaces let you control individual features by corresponding z values directly to them
- There are supervised and unsupervised methods to achieve disentanglement

