




deeplearning.ai


# Image-to-Image Translation

# Outline


- Image-to-image translation
- Other types of translation



# Image-to-Image Translation



Transformation




Available from: <https://twitter.com/citnajt/status/1124904251128406016>

# Image-to-Image Translation



Available from: <https://github.com/NVIDIA/pix2pixHD>


# Image-to-Image Translation



Available from: <https://youtu.be/3RYNThid23g>

# Paired Image-to-Image Translation

Labels to facade




Black-and-white to color



Available from: <https://arxiv.org/abs/1611.07004>

# Paired Image-to-Image Translation

Day to night



Edges to photo




Available from: <https://arxiv.org/abs/1611.07004>

# Paired Image-to-Image Translation



Clothes and pose to  
*pose with clothes*




Available from: <https://arxiv.org/abs/1705.09368>




# Other Translations

“This bird is red with white and has a very short beak”



Available from: <https://arxiv.org/abs/1711.10485>


# Other Translations



Available from: <https://arxiv.org/abs/1905.08233>

# Summary

- Image-to-image translation transforms images into different styles
- GANs' realistic generation abilities are well-suited to image-to-image translation tasks
- Other types of translation include text-to-image or image-to-video






deeplearning.ai

# Pix2Pix Overview

# Outline

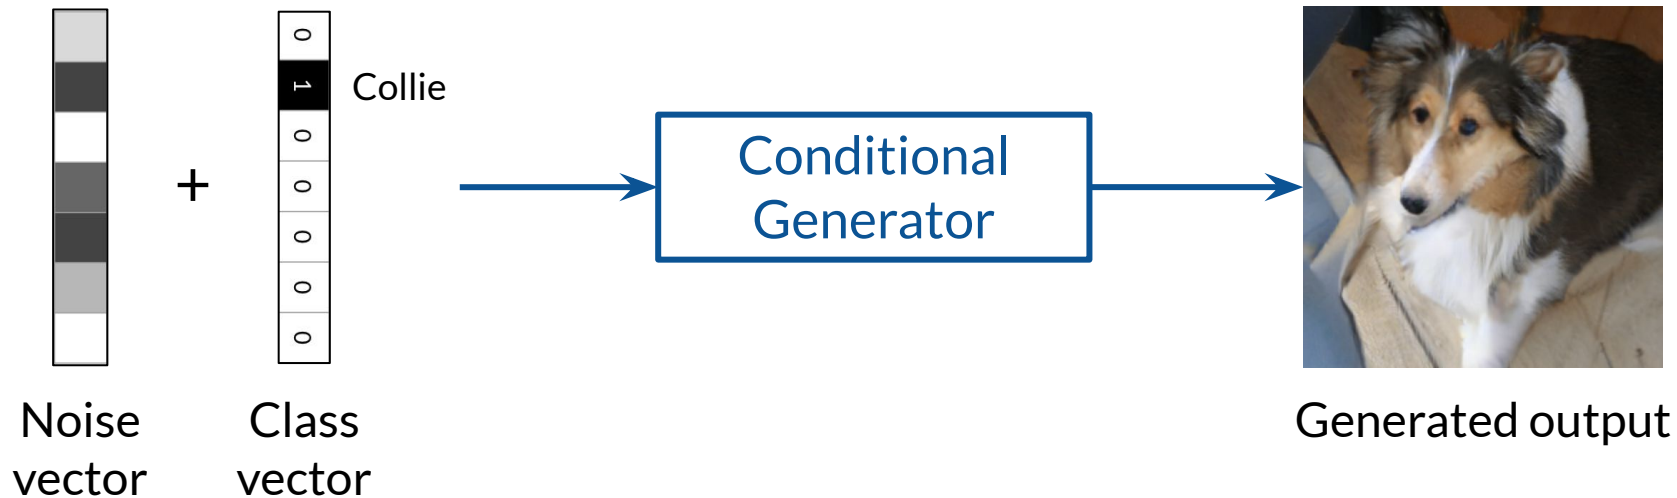
- Overview of Pix2Pix
- Comparison with conditional GAN
- Upgraded generator and discriminator architectures




# Pix2Pix for Paired Image-to-Image Translation

Image-to-Image  $\longrightarrow$  *Pix-to-Pix*  $\longrightarrow$  *Pix2Pix*

# Pix2Pix Generator




# Pix2Pix Generator



Available from: <https://arxiv.org/abs/1611.07004>

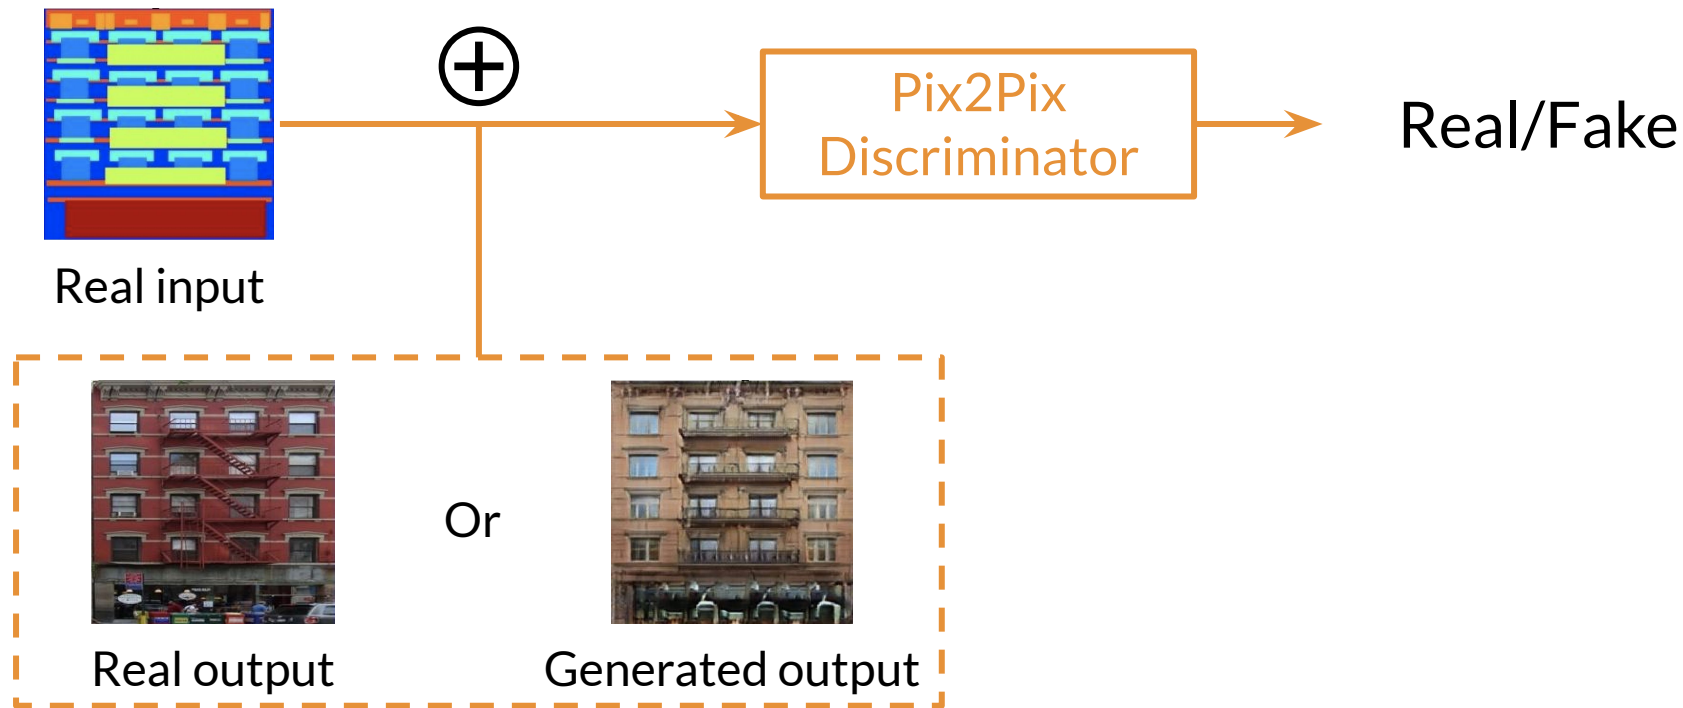


# Pix2Pix Generator




Available from: <https://arxiv.org/abs/1611.07004>

# Pix2Pix Discriminator





Available from: <https://arxiv.org/abs/1611.07004>

# Pix2Pix Upgrades




Based on: <https://arxiv.org/abs/1611.07004>

# Pix2Pix Upgrades




(Left) Based on: <https://arxiv.org/abs/1611.07004>  
(Right) Based on: <https://arxiv.org/abs/1803.07422>

# Pix2Pix Upgrades




Goal is still to produce realistic outputs!



(Left) Based on: <https://arxiv.org/abs/1611.07004>  
(Right) Based on: <https://arxiv.org/abs/1803.07422>

# Summary

- Inputs and outputs of Pix2Pix are similar to a conditional GAN
  - Take in the original image, instead of the class vector
  - No explicit noise as input
- Generator and discriminator models are upgraded






deeplearning.ai

# Pix2Pix: PatchGAN

# Outline

- PatchGAN discriminator architecture
- Matrix output vs. single output





# Pix2Pix Discriminator: PatchGAN




Image available from: <https://arxiv.org/abs/1611.07004>  
Based on: <https://arxiv.org/abs/1803.07422>

# Pix2Pix Discriminator: PatchGAN




Image available from: <https://arxiv.org/abs/1611.07004>  
Based on: <https://arxiv.org/abs/1803.07422>

# Pix2Pix Discriminator: PatchGAN




Image available from: <https://arxiv.org/abs/1611.07004>  
Based on: <https://arxiv.org/abs/1803.07422>

# Pix2Pix Discriminator: PatchGAN





Image available from: <https://arxiv.org/abs/1611.07004>  
Based on: <https://arxiv.org/abs/1803.07422>

# Pix2Pix Discriminator: PatchGAN




Still uses BCE Loss

Image available from: <https://arxiv.org/abs/1611.07004>  
Based on: <https://arxiv.org/abs/1803.07422>

# Pix2Pix Discriminator: PatchGAN

Values closer to *0* = *fake* label




Still uses BCE Loss

Image available from: <https://arxiv.org/abs/1611.07004>  
Based on: <https://arxiv.org/abs/1803.07422>

# Pix2Pix Discriminator: PatchGAN

Values closer to **1** = **real** label




Still uses BCE Loss

Image available from: <https://arxiv.org/abs/1611.07004>  
Based on: <https://arxiv.org/abs/1803.07422>

# Summary

- PatchGAN discriminator outputs a matrix of values, each between 0 and 1
- Label matrices:
  - 0's = fake
  - 1's = real








deeplearning.ai

# Pix2Pix: U-Net

# Outline

- Net framework
  - Encoder-Decoder
- U-Skip connections
- Pix2Pix generator



# Image Segmentation




Image Segmentation

Available from: <https://developer.nvidia.com/blog/image-segmentation-using-digits-5/>

# Image Segmentation




Image-to-Image  
Translation





Image Segmentation


Available from: <https://developer.nvidia.com/blog/image-segmentation-using-digits-5/>

# U-Net Framework: Encoder-Decoder




Based on: <https://arxiv.org/abs/1611.07004>

# U-Net Framework: Skip Connections



Based on: <https://arxiv.org/abs/1611.07004>

# U-Net Framework: Skip Connections



Forward pass


**Skip connections**  
allow information  
flow to the decoder

Based on: <https://arxiv.org/abs/1611.07004>

# U-Net Framework: Skip Connections

Backward pass


**Skip connections**  
improve gradient  
flow to encoder



Based on: <https://arxiv.org/abs/1611.07004>



# Pix2Pix Encoder



# Pix2Pix Encoder

$\chi$  Input size: 256 x 256 x 3




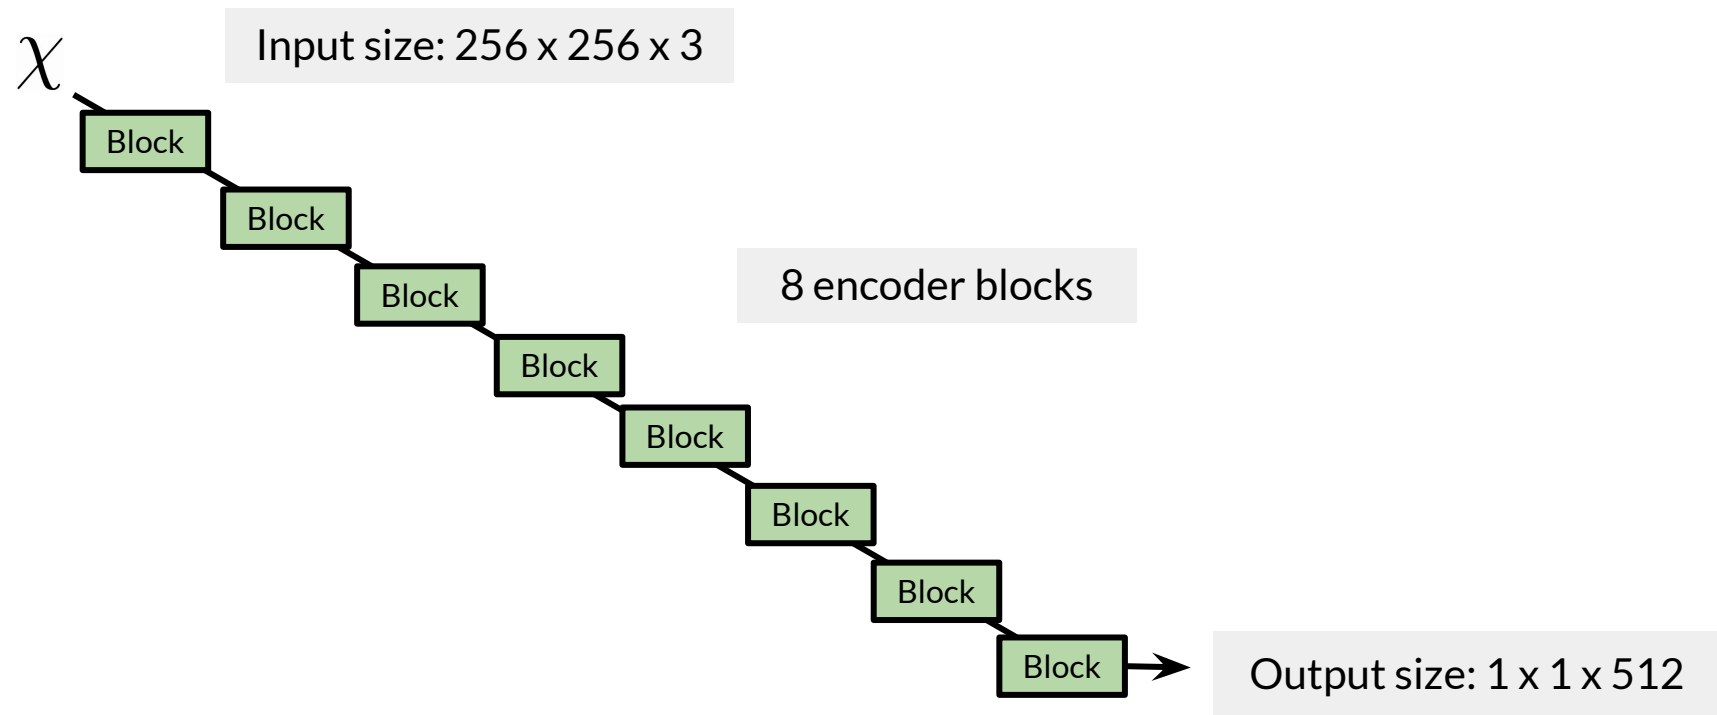


Image available from: <https://arxiv.org/abs/1611.07004>

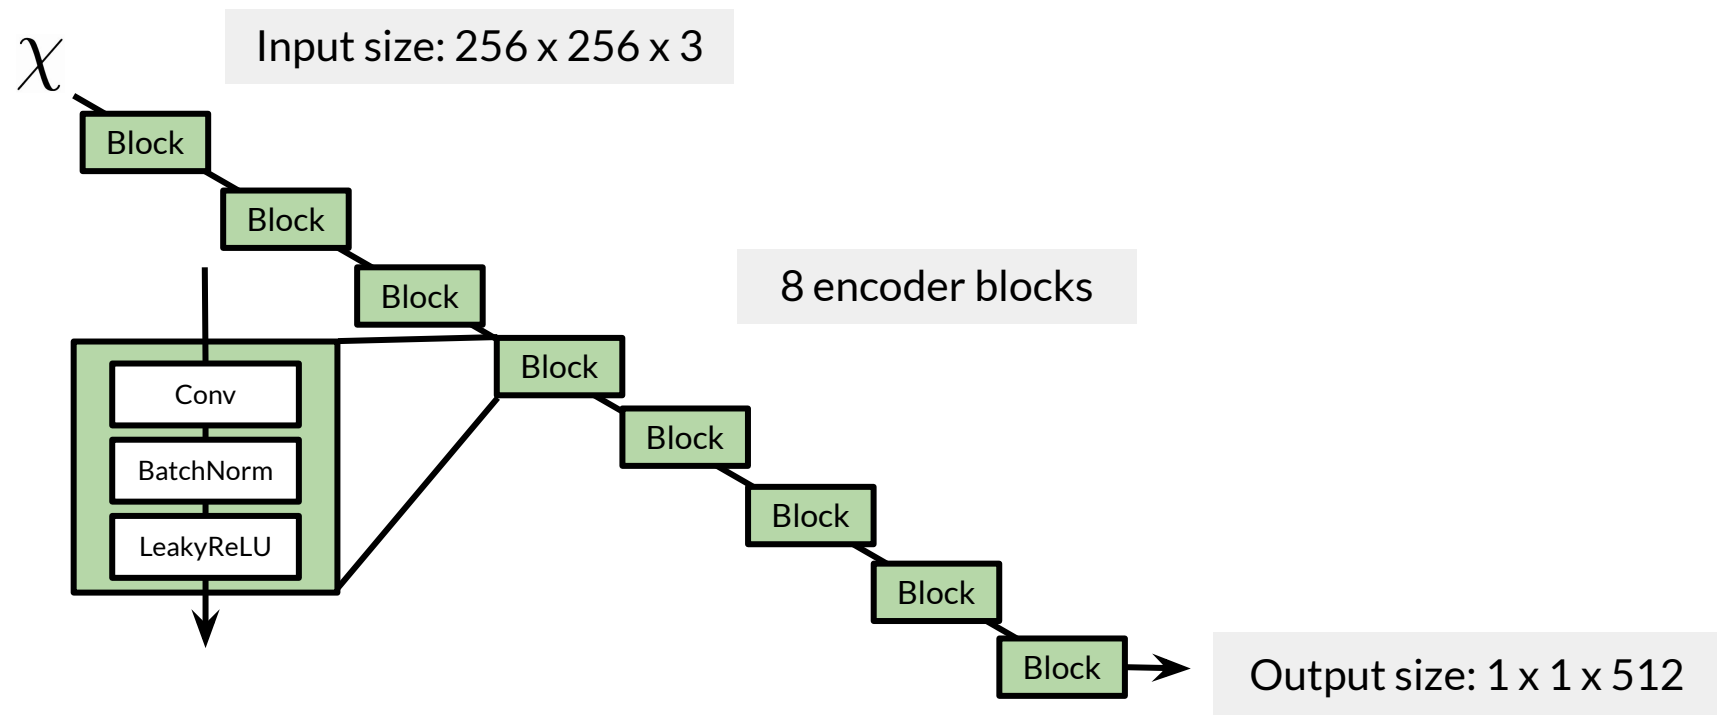
# Pix2Pix Encoder




# Pix2Pix Encoder



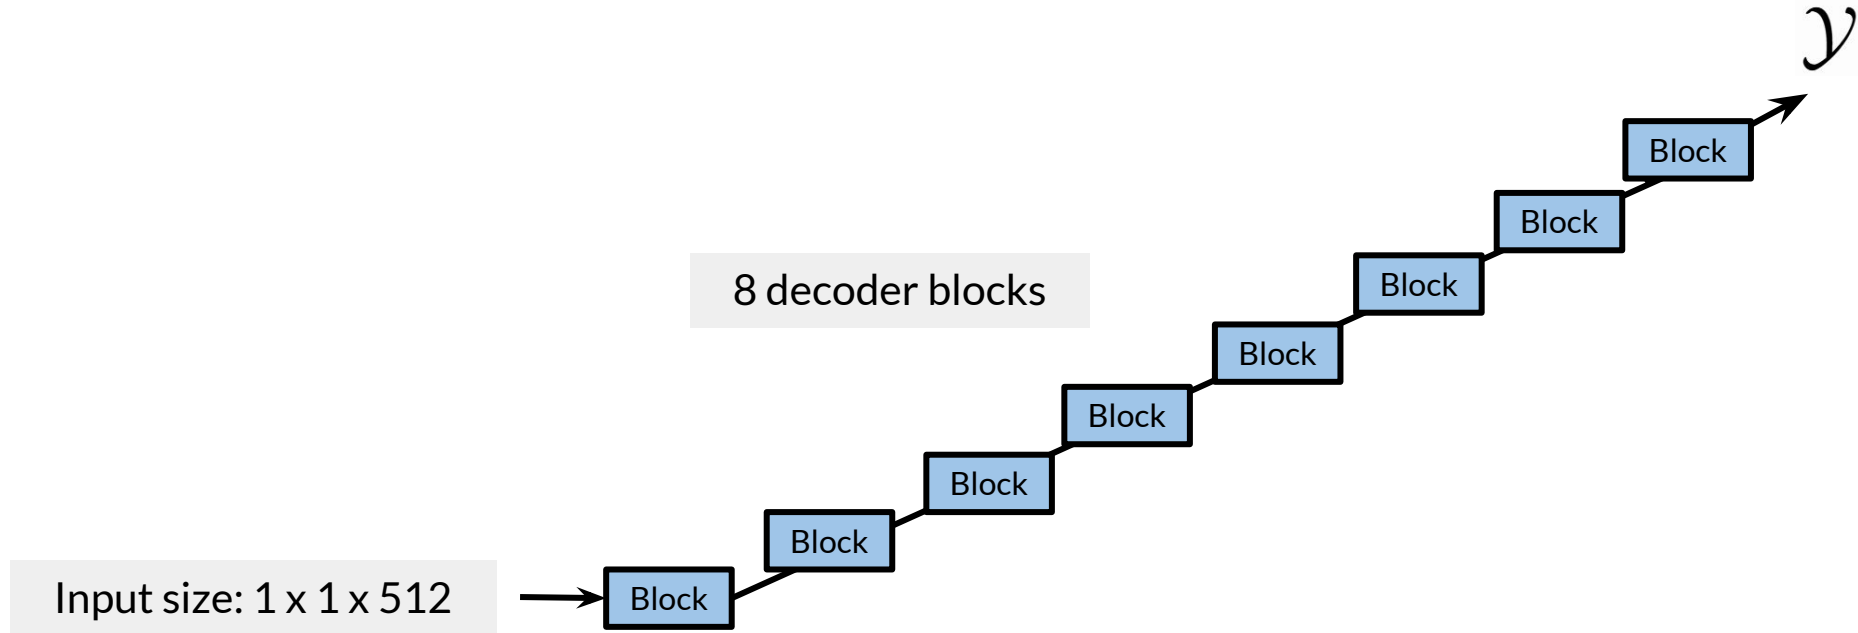
# Pix2Pix Encoder



# Pix2Pix Decoder




# Pix2Pix Decoder



# Pix2Pix Decoder

Output size: 256 x 256 x 3



$y$

8 decoder blocks

Input size: 1 x 1 x 512

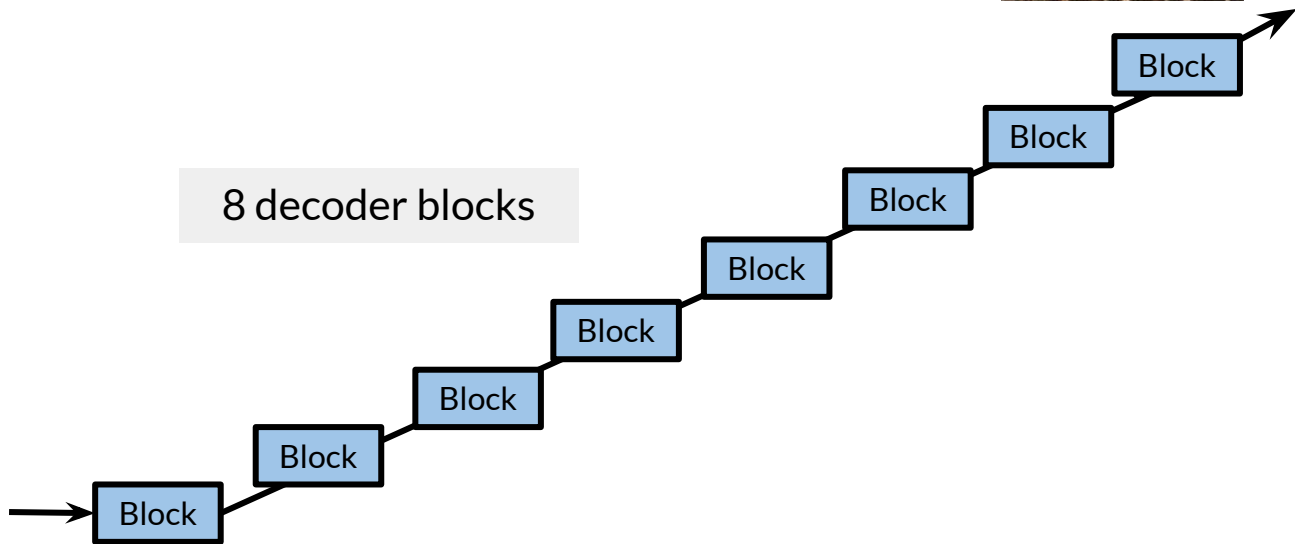


Image available from: <https://arxiv.org/abs/1611.07004>



# Pix2Pix Decoder

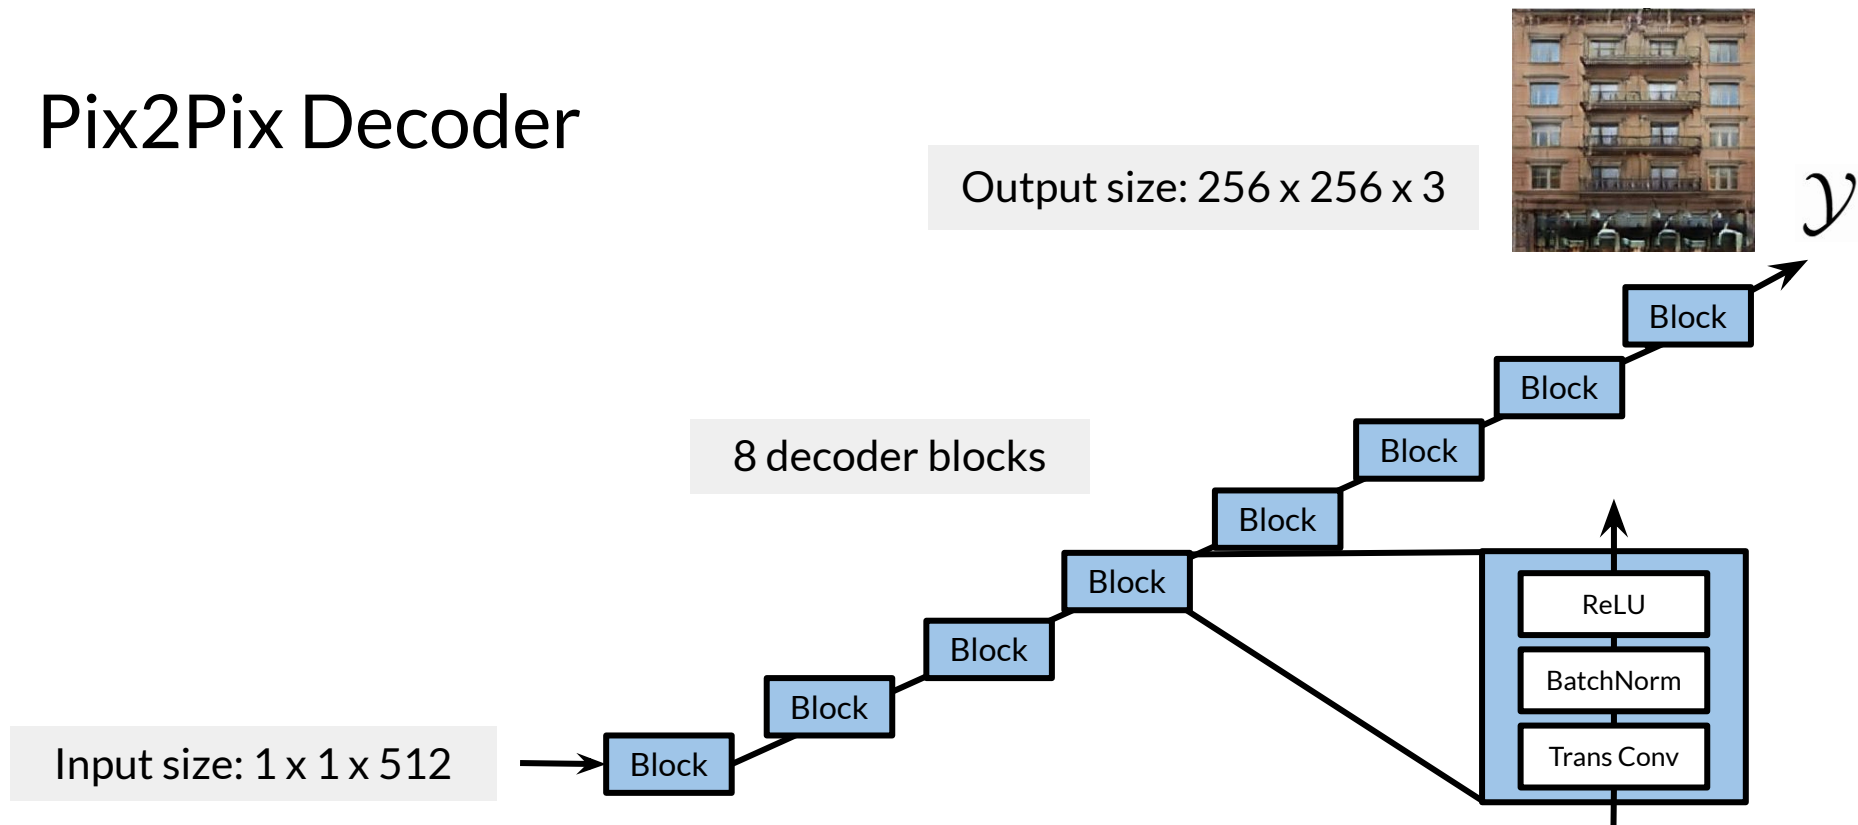



Image available from: <https://arxiv.org/abs/1611.07004>

# Pix2Pix Decoder

Output size: 256 x 256 x 3



$y$

8 decoder blocks

**Dropout** in some decoder blocks  
*adds noise to the network*

Input size: 1 x 1 x 512


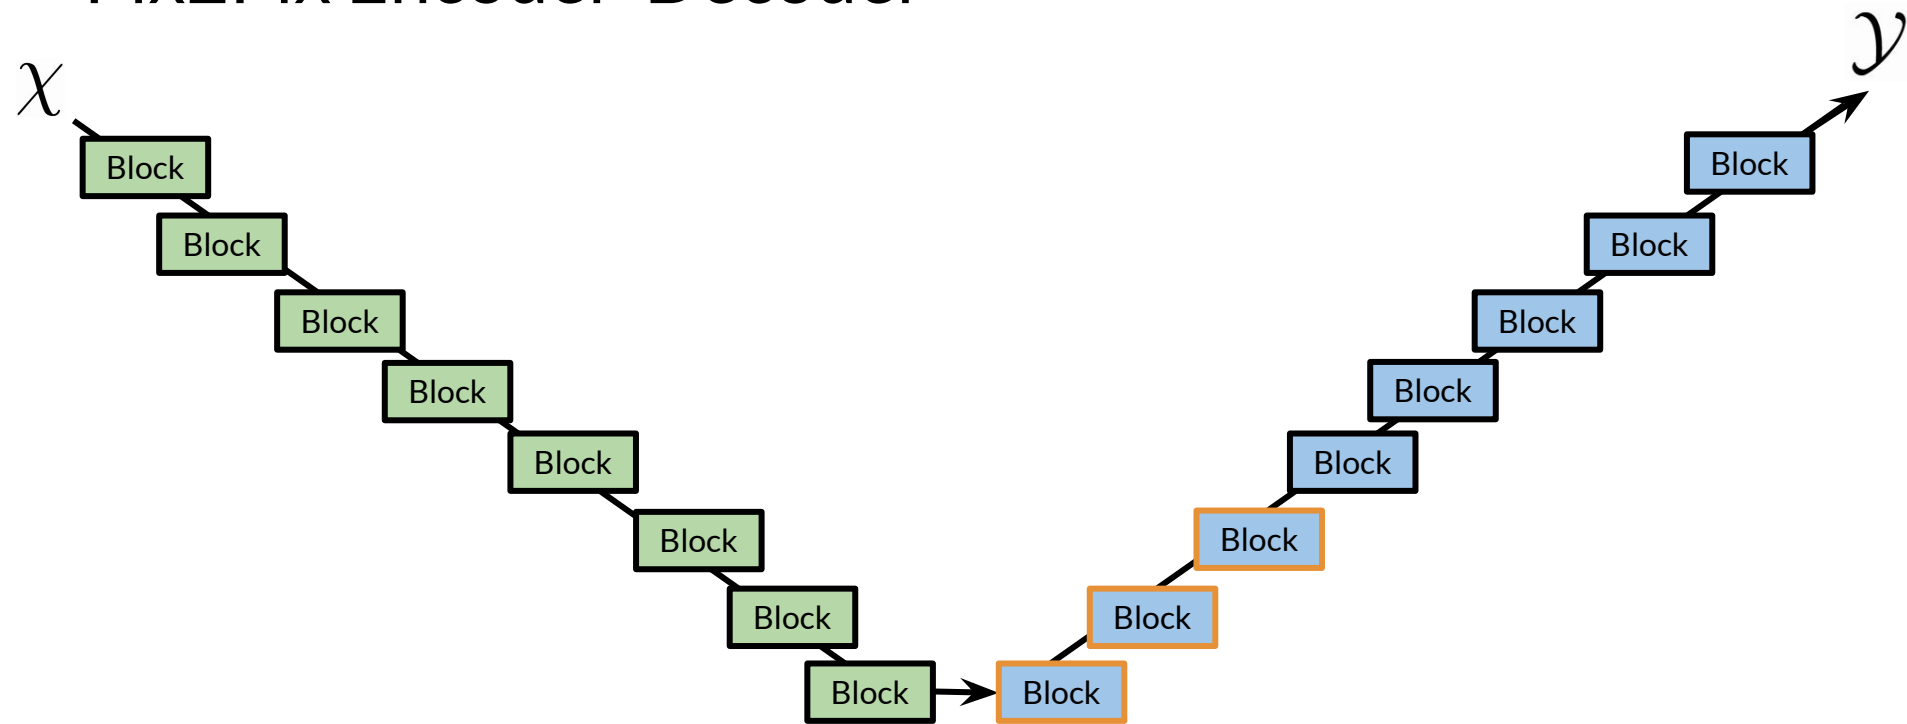


Image available from: <https://arxiv.org/abs/1611.07004>

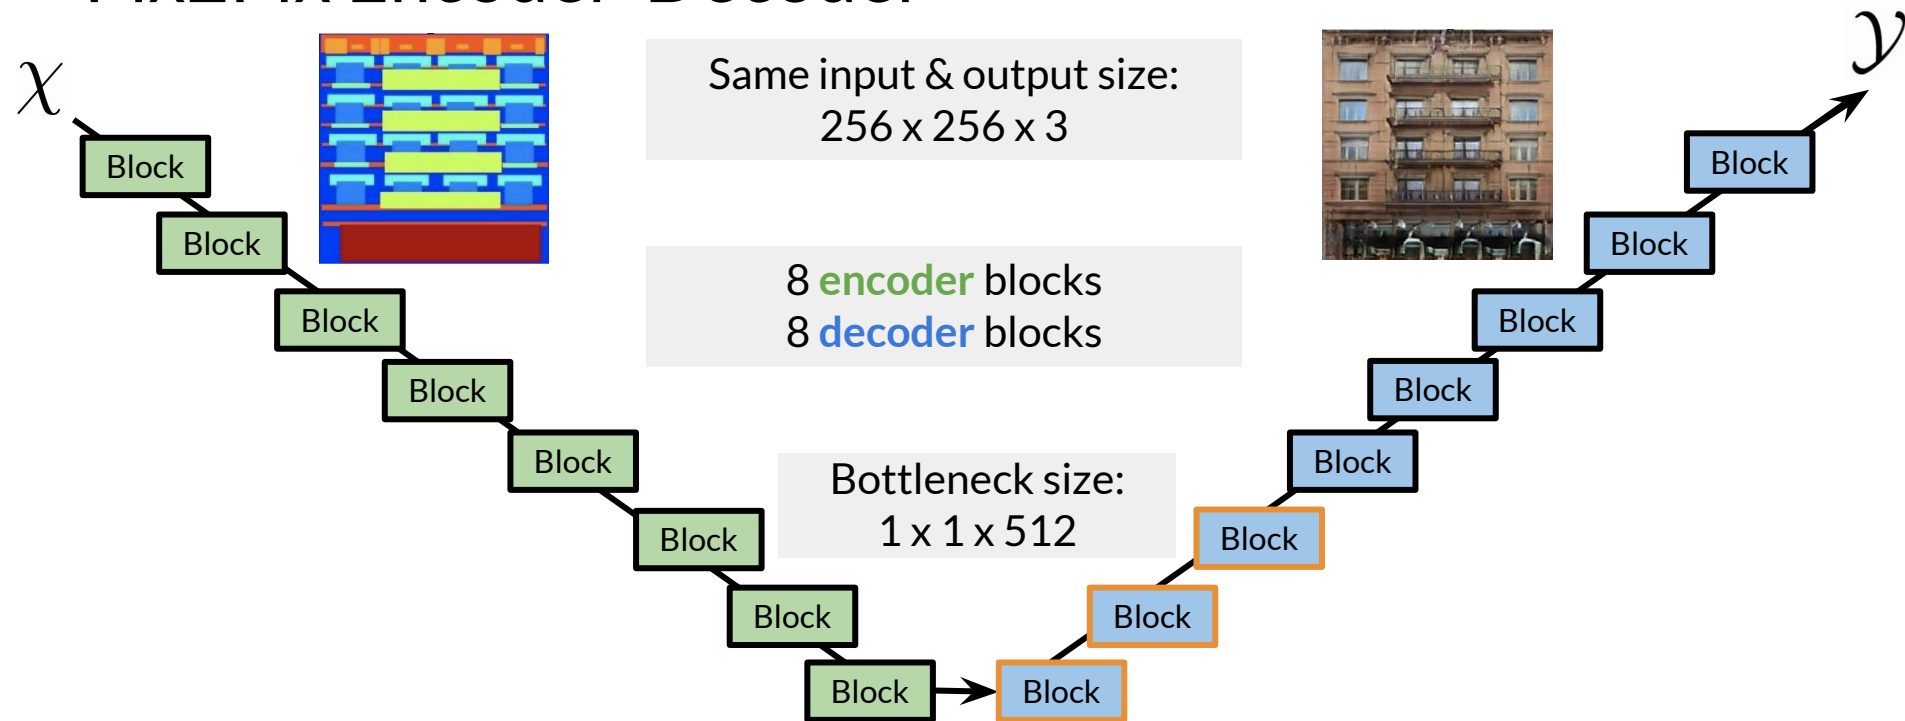
# Pix2Pix Encoder-Decoder





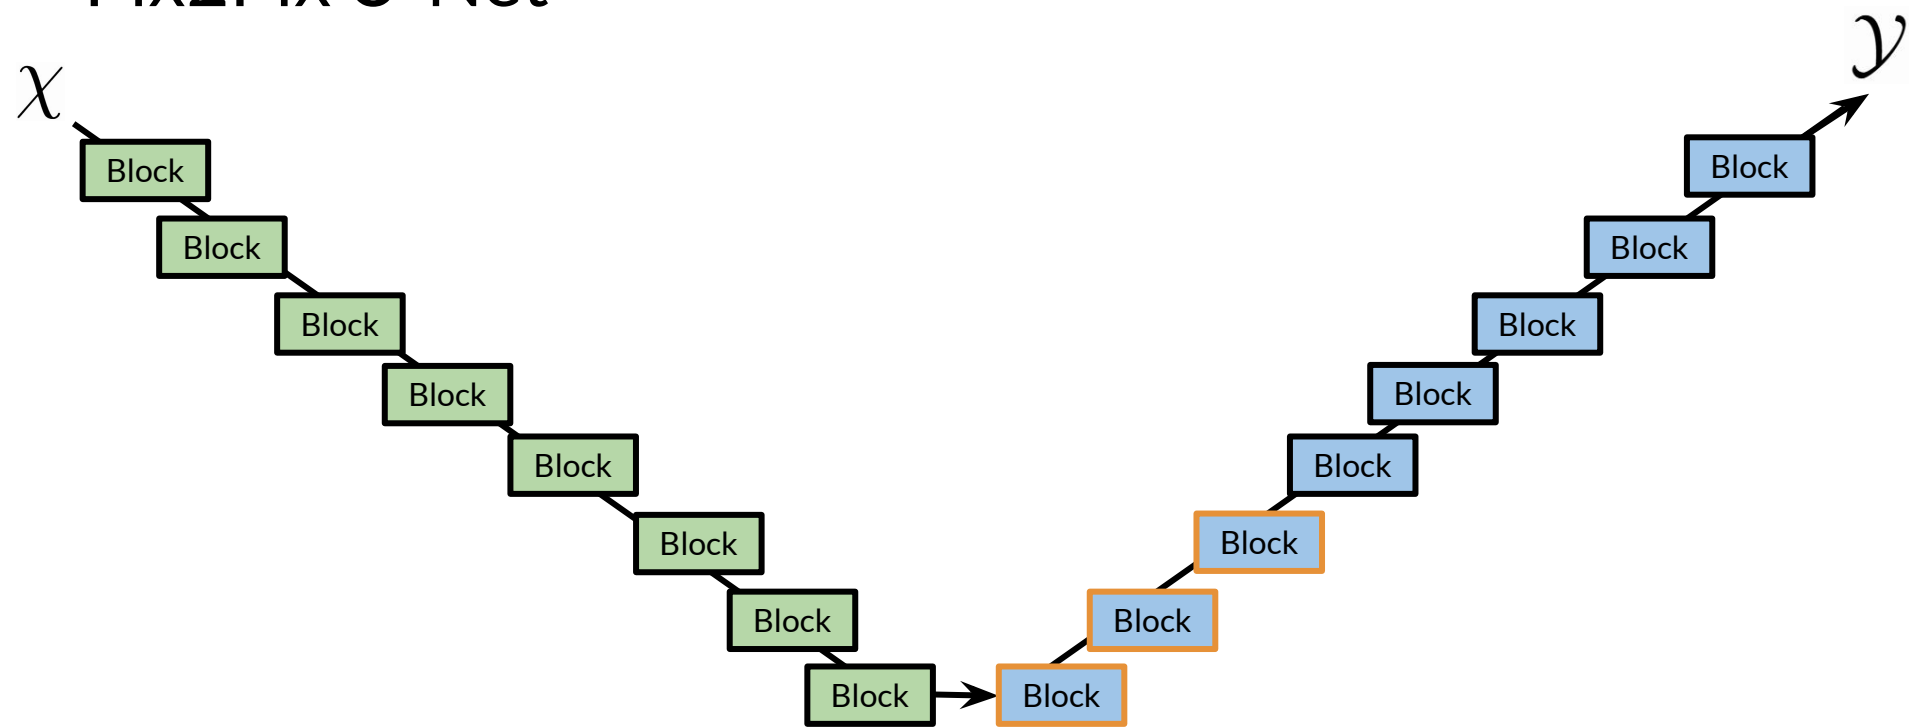


# Pix2Pix Encoder-Decoder

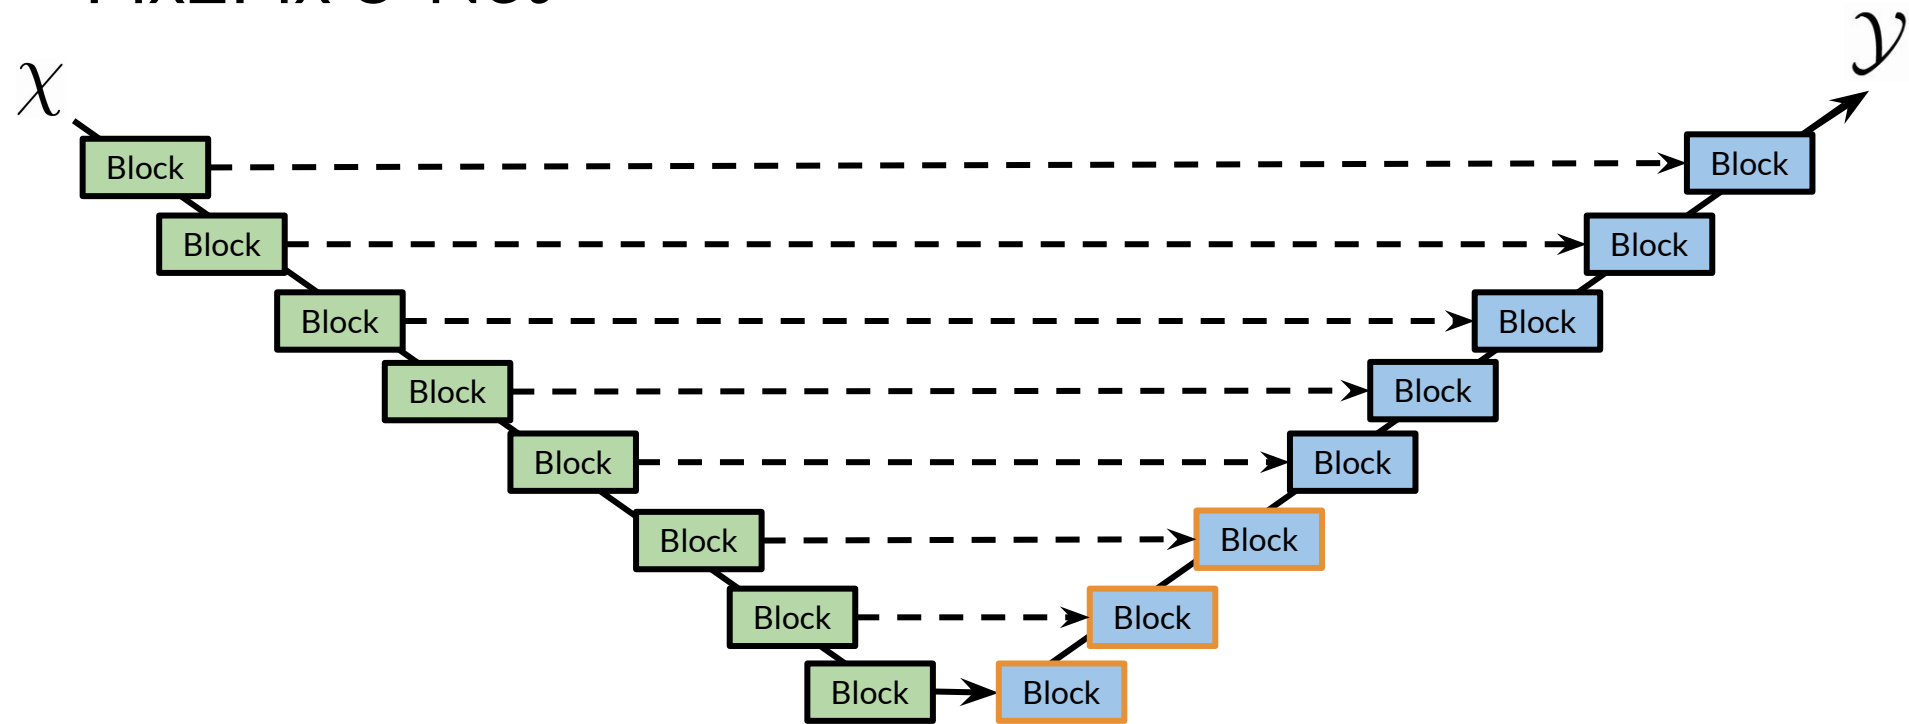


Images available from: <https://arxiv.org/abs/1611.07004>

# Pix2Pix U-Net

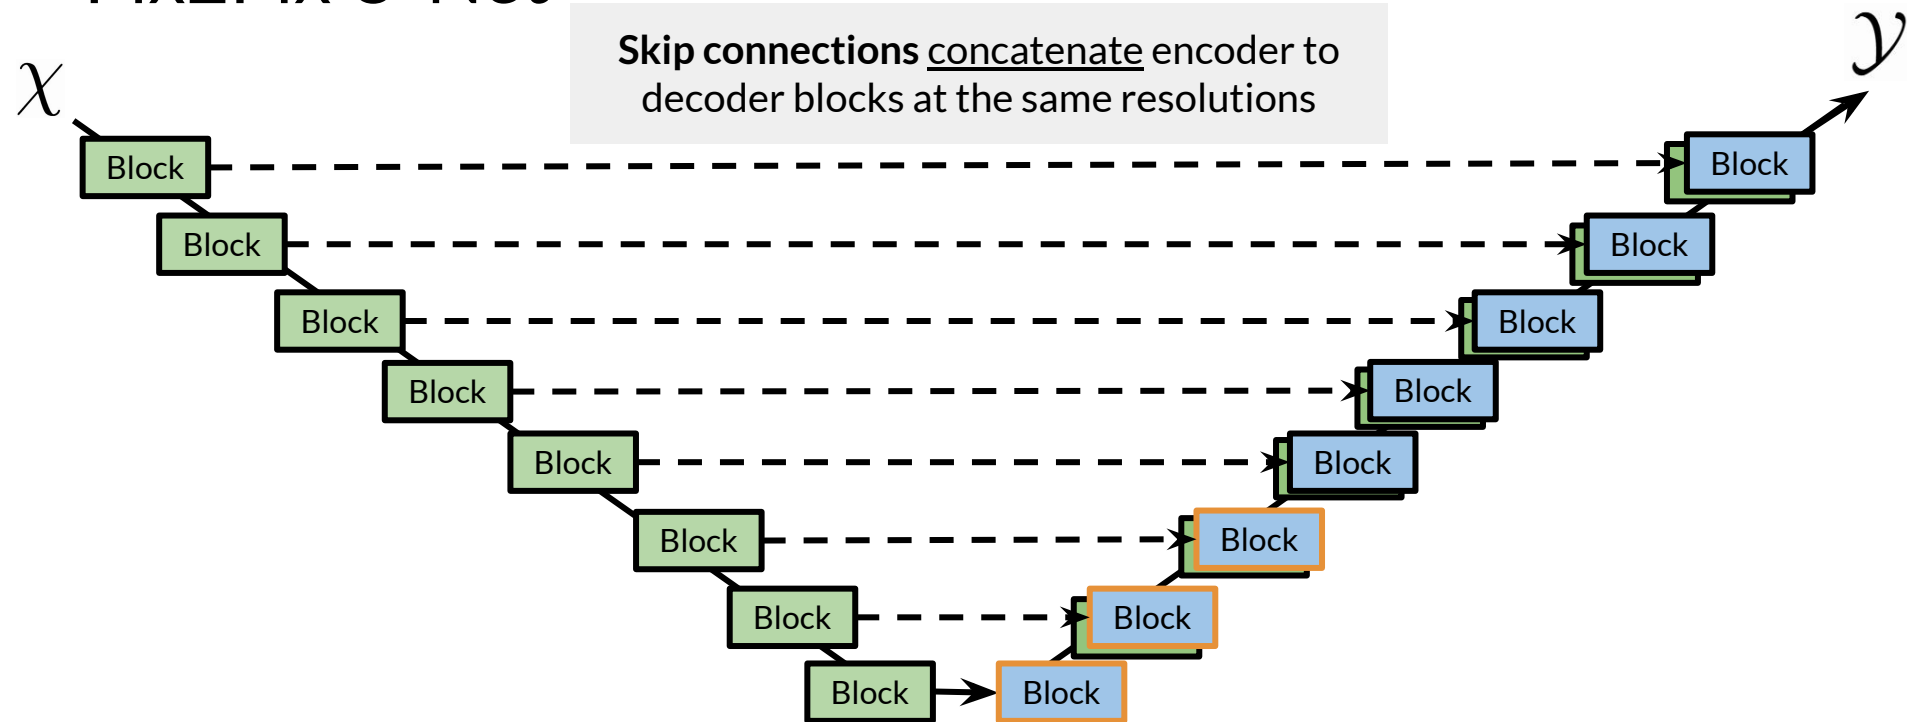


# Pix2Pix U-Net






# Pix2Pix U-Net



# Summary

- Pix2Pix's generator is a U-Net
- U-Net is an encoder-decoder, with same-size inputs and outputs
- U-Net uses skip connections
  - Skip connections help the decoder learn details from the encoder directly
  - Skip connections the encoder learn from more gradients flowing from decoder






deeplearning.ai

# Pix2Pix: Pixel Distance Loss Term

# Outline

- Regularization and additional loss term
- Encourage pixel distance between generated and real outputs
- Additional loss term for Pix2Pix generator



# Additional Loss Term

$$\min_g \max_c \text{Adversarial Loss} + \lambda * \text{Other loss term}$$

# Additional Loss Term

$$\min_g \max_c \text{Adversarial Loss} + \lambda * \text{Pixel loss term}$$

# Pixel Distance Loss Term

$$\sum_{i=1}^n$$



Generated output

—



Real output

# Pix2Pix Generator Loss

BCE Loss +

$$\lambda \sum_{i=1}^n$$



-



Available from: <https://arxiv.org/abs/1611.07004>



# Pix2Pix Generator Loss

$$\text{BCE Loss} + \lambda \sum_{i=1}^n \left| \text{generated\_output} - \text{real\_output} \right|$$

# Summary

- Pix2Pix adds a Pixel Distance Loss term to the generator loss function
- This loss term calculates the difference between the fake and the real target outputs
- Softly encourages the generator with this additional supervision
  - The target output labels are the supervision
  - Generator essentially “sees” these labels






deeplearning.ai


# Pix2Pix: Putting It All Together

# Outline

- Put the Pix2Pix architecture together!
  - U-Net generator
  - Pixel Distance Loss term
  - PatchGAN discriminator




# Pix2Pix



Real input


Image available from: <https://arxiv.org/abs/1611.07004>

# Pix2Pix



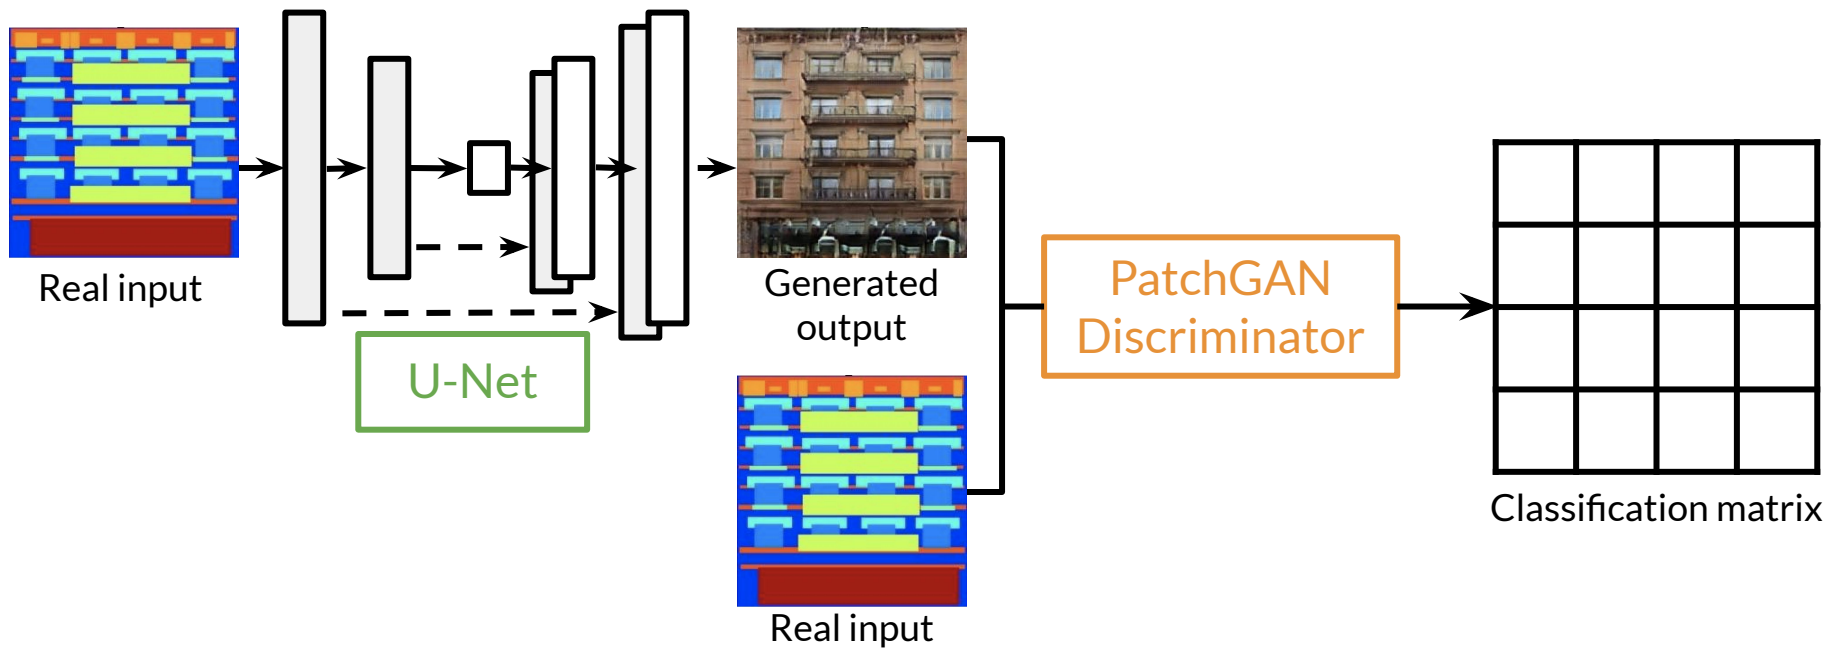
Images available from: <https://arxiv.org/abs/1611.07004>

# Pix2Pix



Images available from: <https://arxiv.org/abs/1611.07004>

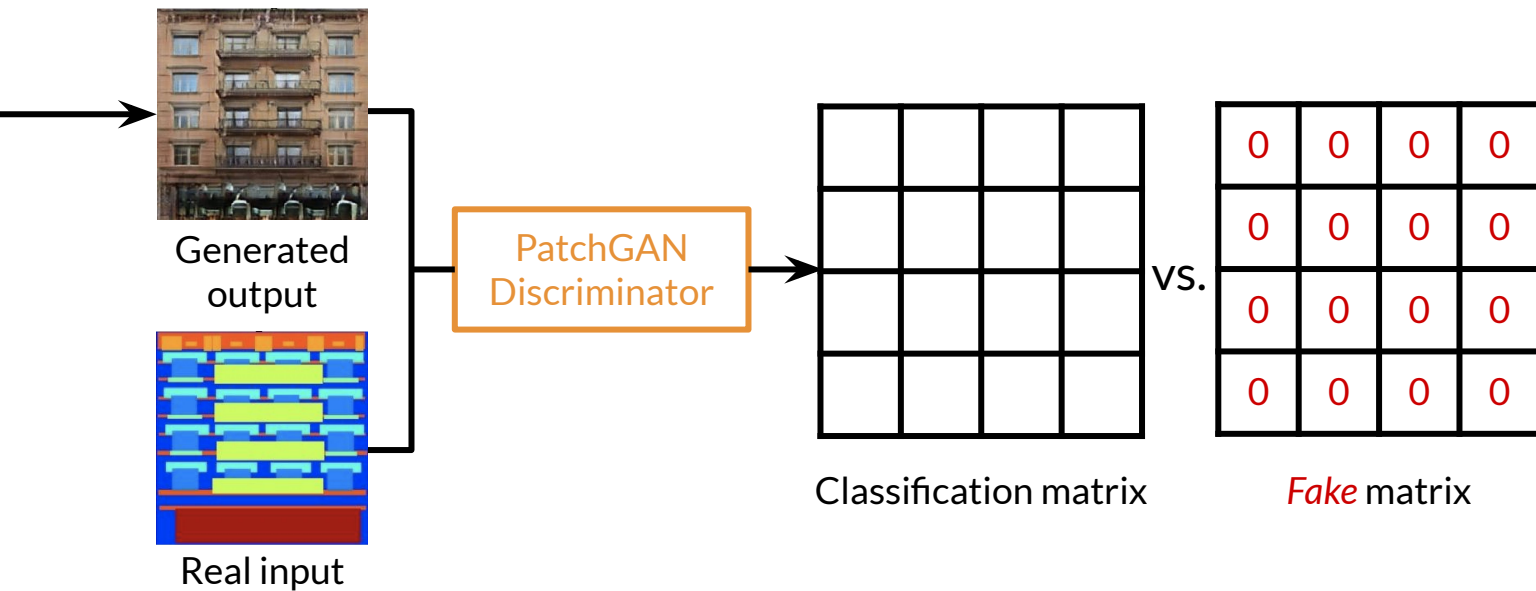
# Pix2Pix



Images available from: <https://arxiv.org/abs/1611.07004>

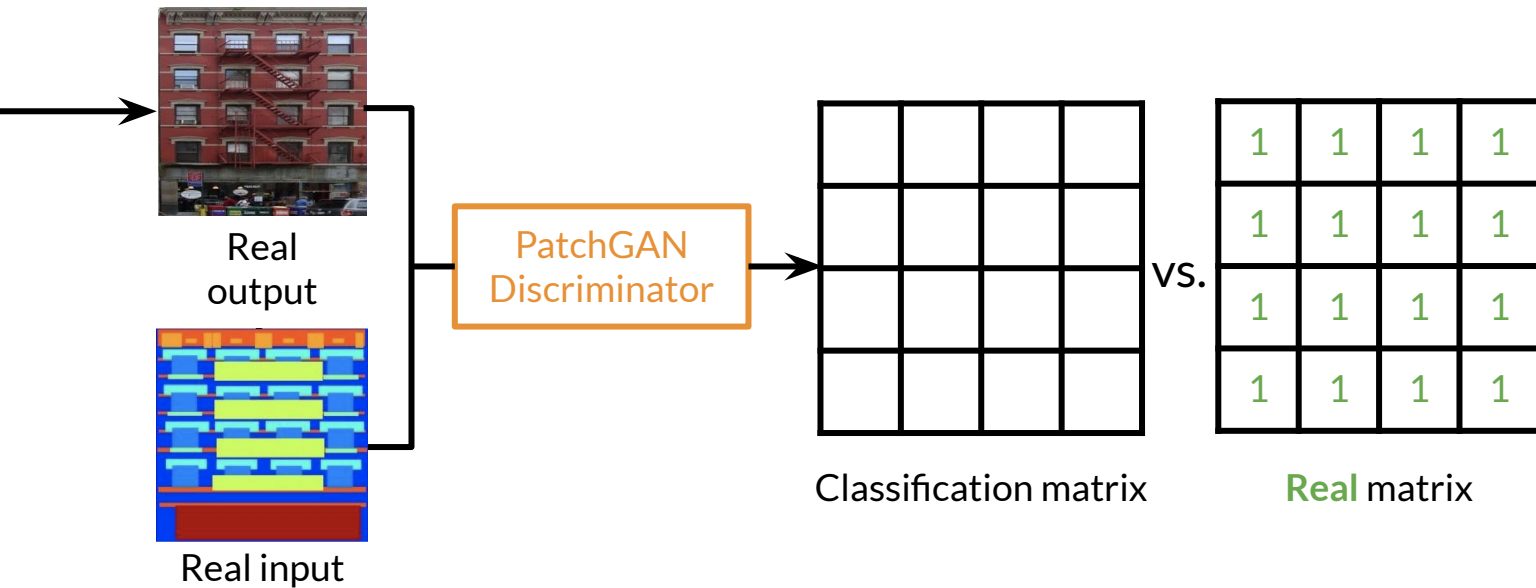


# Pix2Pix: Discriminator Loss



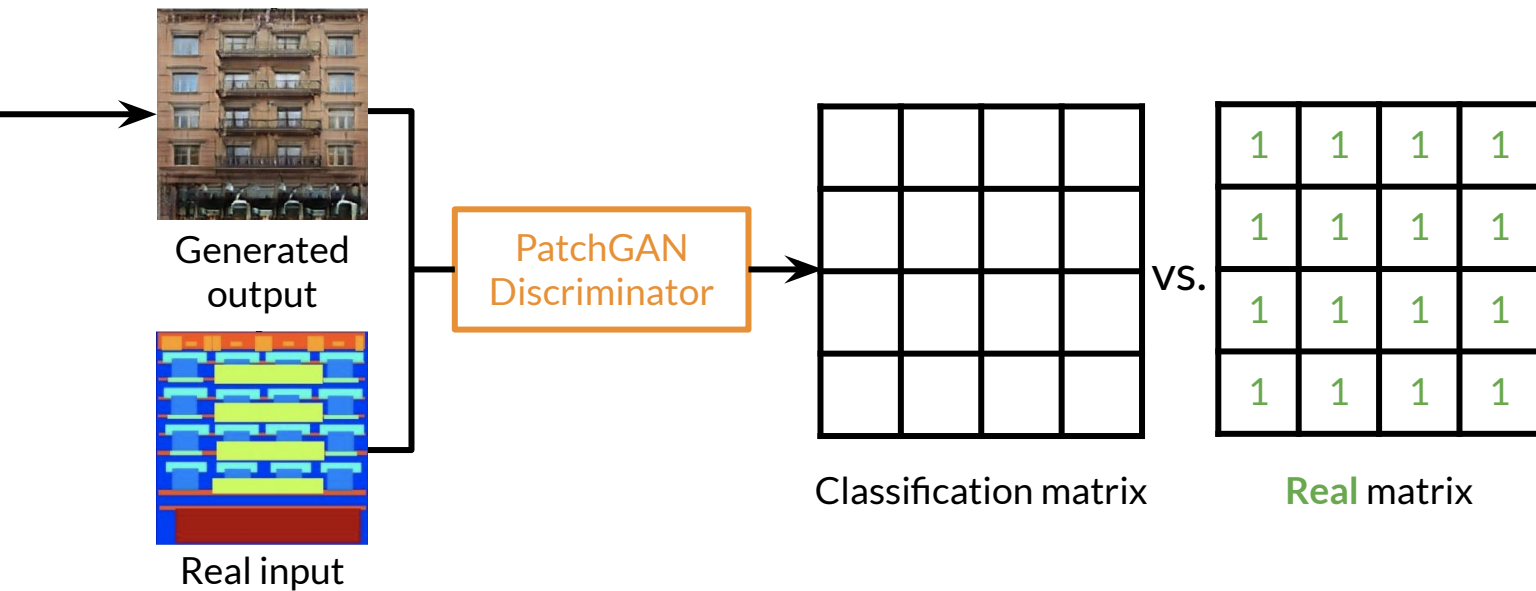
Images available from: <https://arxiv.org/abs/1611.07004>

# Pix2Pix: Discriminator Loss



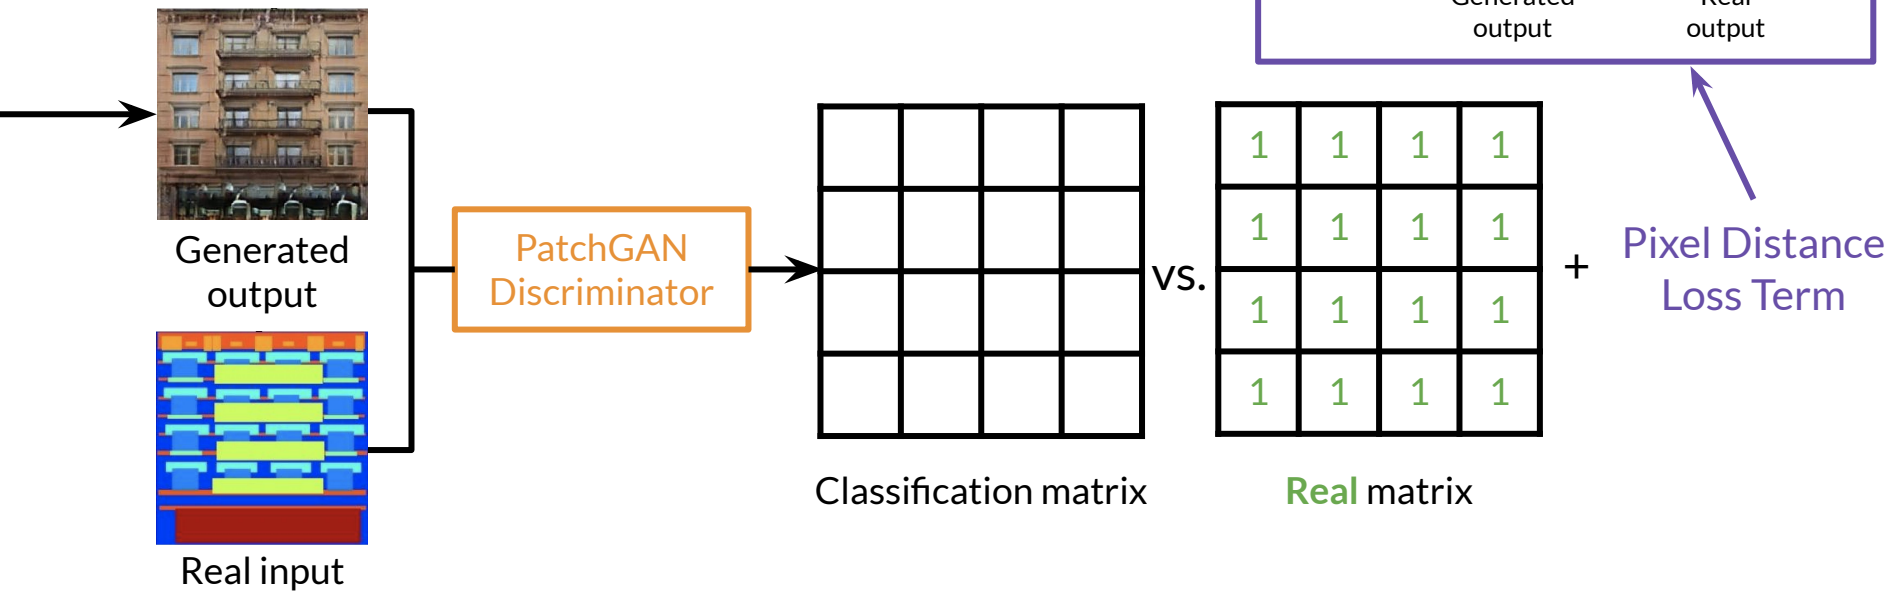
Images available from: <https://arxiv.org/abs/1611.07004>

# Pix2Pix: Generator Loss



Images available from: <https://arxiv.org/abs/1611.07004>


# Pix2Pix: Generator Loss



Images available from: <https://arxiv.org/abs/1611.07004>

# Summary

- U-Net generator: image  $\rightarrow$  image
- PatchGAN discriminator
  - Inputs input image and paired output (either real target or fake)
  - Outputs classification matrix
- Generator loss has a regularization term





deeplearning.ai

# Pix2Pix Advancements

# Outline

- Improvements and extensions of Pix2Pix for paired image-to-image translation
  - Higher resolution images
  - Image editing




# Pix2PixHD



Available from: <https://github.com/NVIDIA/pix2pixHD>



# GauGAN



Available from: <https://blogs.nvidia.com/blog/2019/03/18/gaugan-photorealistic-landscapes-nvidia-research/>

# Summary

- Pix2PixHD and GauGAN are successors of Pix2Pix
- They are designed for higher resolution images
- They highlight opportunities for image editing using paired image-to-image translation
  - Pix2Pix can do this too, of course!

